

Electromagnetic Momentum Engineering

How Fields Push, Steer, and Structure the Physical World

Dustin Lee

Introduction

This book is about a tool that's been sitting in plain sight for over a hundred years. Electromagnetic fields carry momentum. That momentum can be transferred to matter. The transfer can be engineered deliberately. And yet, outside of a handful of exotic applications, laser cooling, solar sails, optical tweezers, almost nobody treats this as a general engineering tool. That's a gap worth closing.

The broader framework this book sits inside is called Unified Substrate Theory (UST) and Universal Baseline Level (UBL). The core idea is that the physical substrate, the medium through which all fields propagate, is itself an engineering target. EM momentum engineering, in this context, is the discipline of deliberately shaping, compressing, redirecting, and harnessing electromagnetic fields to do useful work at a structural level. Not just to carry signals. To push things.

If you've ever wondered why radiation pressure is covered in one paragraph in a graduate EM course and then never mentioned again, this book is the answer. It shouldn't be a footnote. It should be a field.

What EM Momentum Actually Is

The Basic Fact Nobody Emphasizes

Electromagnetic fields carry momentum. This isn't a theoretical curiosity or a quantum weirdness thing. It's classical electrodynamics, fully contained in Maxwell's equations, and it's been experimentally confirmed since 1903. The momentum is real. It can be transferred to objects. Those objects move.

The momentum density of an electromagnetic field is given by the Poynting vector divided by the square of the speed of light. The Poynting vector itself is the cross product of the electric field E and the magnetic field B , divided by the permeability of free space. Written out: $\mathbf{g} = \mathbf{S}/c^2 = (\mathbf{E} \times \mathbf{B}) / (\mu_0 c^2)$. That's momentum per unit volume, sitting in the field, right now, everywhere there's a field.

If you integrate that over a volume, you get total field momentum. If that field dumps its momentum into a surface, an absorbing wall, a charged particle, a dielectric interface, you get a force. The force per unit area on an absorbing surface is the field's energy density divided by the speed of light. For a reflecting surface, it's twice that, because the momentum reverses direction and you get both the original push and the reflected push.

The Maxwell Stress Tensor

The Maxwell stress tensor is the machinery that tells you exactly what force the field exerts on any surface, in any direction, at any point. It's a 3×3 matrix. Each component gives the flux of one component of momentum in one direction. Integrate the tensor over a closed surface surrounding an object, and you get the total electromagnetic force on that object.

The tensor isn't exotic. It's just the bookkeeping tool for momentum conservation in the field. You use it the same way you use the stress tensor in structural mechanics, to figure out where the force is coming from and where it's going.

The key insight: the stress tensor components involve squared field magnitudes. That means force scales with field intensity squared. Double the field amplitude, quadruple the force. This is why resonant cavities

and coherent superposition are such powerful tools, you're not just adding field strength linearly, you're adding it in a way that multiplies force nonlinearly.

Photons and Momentum Without Mass

Photons carry momentum even though they have zero rest mass. The momentum of a photon is $\mathbf{p} = \mathbf{E}/c = \hbar/\lambda$, where E is its energy, c is the speed of light, $\hbar$ is Planck's constant, and λ is wavelength. This is relativistically consistent and you don't need mass to have momentum, you need energy and motion. Photons have both.

In classical terms, this is the same thing. The field has energy density. The energy density moves at c . That moving energy carries momentum. The photon picture and the field picture are two descriptions of the same physics.

Why This Gets Ignored

Radiation pressure shows up in engineering curricula in exactly the following contexts: solar sails, laser cooling of atoms, and optical tweezers. Then it disappears. The reason is scale. At room-temperature power levels and normal engineering geometries, radiation pressure is tiny compared to mechanical forces. A kilowatt of light on a perfect reflector produces about 6.7 micronewtons. A kilowatt of airflow produces orders of magnitude more force.

So the engineering culture decided radiation pressure was interesting but not useful, and moved on.

That's a mistake, and it's the mistake this book is correcting. The forces are small only when the field geometry is unoptimized. When you compress a field into a resonant cavity, concentrate it with high-permittivity materials, and phase-stack multiple sources, the force per unit power goes up dramatically. More importantly, when the application doesn't need large forces, precision positioning, drag compensation, contactless manipulation, radiation-pressure-scale forces are exactly right.

The Momentum Is Always There

Every antenna you've ever built radiates EM momentum. Every microwave oven is full of it. Every RF cavity in a particle accelerator has it bouncing around inside. The field's momentum doesn't care whether you're thinking about it. It's exerting forces on those cavity walls right now. The question isn't whether it's there, it's whether you're designing to use it.

CHAPTER TWO

The Substrate Perspective: Fields Aren't Empty

The Standard Story

Standard physics treats the vacuum as featureless. Fields propagate through it at c . There's no medium. No structure. No geometry. Maxwell's equations are written in a vacuum, and the vacuum is just empty spacetime with defined constants, the permittivity ϵ_0 and the permeability μ_0 , that don't vary, don't have structure, and don't push back.

That picture works for most purposes. It gives you antenna theory, circuit theory, and waveguide design, and those work fine. But it's an engineering approximation, not a complete description, and it starts to break down at the edges of what you're trying to do.

The Substrate Hypothesis

The Unified Substrate Theory starts from a different premise: the vacuum isn't empty. It's a medium. It has geometry. It has local density. It has resistance to change. The constants ϵ_0 and μ_0 aren't arbitrary numbers, they're properties of the substrate. The speed of light $c = 1/\sqrt{\epsilon_0\mu_0}$ isn't just a coincidence of unit choices. It's the propagation speed of a disturbance through a medium with those properties.

This matters for engineering because a medium can be shaped. A medium can be biased. A medium can be engineered locally if you know what affects it. In the standard picture, you have no leverage on the vacuum. In the substrate picture, local conditions, mass concentrations, field gradients, moving charges, warp the substrate geometry, and that warp affects how fields propagate and what forces they exert.

Think of It as a Lattice

The easiest mental model for the substrate isn't atoms. It isn't a classical ether in the 19th-century sense either. Think of it as a lattice, a structured medium at a level below atoms, where each lattice point has some compliance, some inertia, and some restoring force. An EM wave is a disturbance propagating through that lattice. The lattice transmits it, but the lattice also pushes back.

In a uniform lattice, the field propagates cleanly at c , and everything looks like the standard vacuum model. But if the lattice geometry is perturbed, compressed in one region, expanded in another, anisotropic along one axis, then propagation changes. The field slows down, bends, intensifies, or redistributes its momentum. These are real, observable effects in the standard framework too: gravitational lensing, frame dragging, the Shapiro delay. UST treats those as substrate geometry effects made large by extreme mass concentrations.

At engineering scales, mass concentrations don't do much to the substrate. But field concentrations can. And that's the lever.

Why Vacuum Structure Matters for Force

Here's the engineering implication. In the standard model, EM force on an object comes exclusively from field interaction with charged particles in that object. The field hits the charges, the charges accelerate, momentum transfers. Clean and simple.

In the substrate model, there's an additional channel: the field can interact with the substrate geometry itself. The reaction force of that interaction doesn't propagate through photon recoil, it propagates through the substrate. This changes the accounting. It means that in principle, a field can exert a force on a system in a way that doesn't look like conventional radiation pressure, because the momentum is being transferred to and through the substrate, not just to the charged surfaces of the cavity.

This is speculative at the engineering level. It's theoretically motivated. And it's the reason the open problems in Chapter 9 are interesting rather than closed.

What's Not Speculative

The substrate picture doesn't change the established physics. Radiation pressure is real either way. Maxwell's stress tensor works either way. Resonant cavity compression is real either way. What the substrate perspective adds is a richer engineering vocabulary: instead of just asking "what force does this field exert on these charges," you ask "what is the local geometry, and how does the field interact with that geometry?" That question has more handles.

CHAPTER THREE

Lattice Geometry: The Shape of the Field's Home

Geometry Isn't Uniform

Whether you're thinking about the physical substrate in the UST sense or just the local electromagnetic environment in the standard sense, one thing is clear: the space a field propagates through isn't uniform. There are regions of high permittivity and low permittivity. Regions where the boundary conditions are rigid (metal walls) and regions where they're soft (open space). Regions where multiple fields are stacked in phase and regions where they cancel. All these affect what the field does and where its momentum goes.

Calling this "lattice geometry" is a useful shorthand. It's not claiming there's a literal crystal lattice in the vacuum. It's saying that the effective propagation environment has structure, and that structure is engineerable.

Compressed vs. Expanded Lattice Regions

When an EM wave travels through a region of higher permittivity, a dielectric, for example, it slows down. The wavelength shortens. The field intensity, for the same power, is compressed into a smaller spatial volume. This is a "compressed lattice region" in the language of substrate geometry. More field energy per unit volume means higher momentum density. Higher momentum density means higher radiation pressure on any surface in that region.

When the wave exits the high-permittivity region, it speeds back up. The field expands. Momentum density drops. At the interface between compressed and expanded regions, there's a momentum gradient, and momentum gradients produce forces. This is not abstract. This is how dielectric lenses work. This is how impedance transformers work. This is how field-matter interaction happens at every material boundary.

A Concrete Example: The Tapered Waveguide

Take a rectangular waveguide. As the waveguide narrows, the field is forced into a smaller cross-sectional area. Field intensity rises. Momentum density rises. At the end wall of a tapered guide, or at the opening of a horn antenna, there's a definite force on that boundary due to radiation pressure.

This is measured. It's routine. Horn antenna designers know the field distribution at the aperture, and the radiation pressure at the aperture walls is part of the mechanical loading on the structure. It's usually tiny and ignored for structural purposes. But it's real, and it's the direct product of field geometry.

Now extend the concept: what if you designed the taper not just for impedance matching or beam pattern, but specifically to maximize the momentum gradient at the end? What if the end of the taper was terminated not in an aperture but in a load, a material object you want to push? You've just designed a radiation pressure actuator.

Anisotropic Lattice Regions

A medium doesn't have to be isotropic. Anisotropic materials, crystals, layered structures, fiber composites, have different properties in different directions. The same is true of EM environments. A waveguide propagates in one direction and not others. A resonant cavity supports certain modes and suppresses others. A metamaterial structure can be designed with an effective index that's positive in one axis and negative in another.

When the effective EM medium is anisotropic, field propagation is directionally biased. Momentum flows preferentially in directions where propagation is favored and is suppressed in directions where propagation is penalized. This is the geometric basis for directional force generation. You're not hoping the field happens to push in the right direction, you're designing the geometry, so it has no choice.

Gradients Are the Mechanism

Every time there's a gradient in field energy density, there's a potential force. The force on a volume element in a field is proportional to the gradient of the field's energy density, weighted by the relevant material properties (permittivity gradient for electric fields, permeability gradient for magnetic fields). This is straightforward from the Maxwell stress tensor.

So the engineering problem reduces to this: create a steep, directed gradient in field energy density. The steeper the gradient, the stronger the force. The more directed it is, the more of that force points usefully. Everything in the next several chapters is about how to create those gradients on purpose.

Practical Implication for Design

When you're designing a system to exploit EM momentum, think geometrically first. Where is the field most intense? Where does it transition from high-density to low-density? Where are the sharp impedance steps? Those are the force sites. Everything else is detail. Get the geometry right, and the forces follow.

CHAPTER FOUR

EM Compression: Packing More Momentum Into Less Space

What Compression Means Here

EM compression is the deliberate act of increasing local field energy density beyond what you'd get from a simple propagating wave. You're taking the field's energy and squeezing it into a smaller volume, which raises the momentum density in that region, which increases the force that region can exert on a surface or a material.

There are three primary tools for doing this: resonant cavities, coherent superposition from phased arrays, and geometry-driven field concentration. They can be combined, and the best designs use all three.

Resonant Cavities as Compression Engines

A resonant cavity is a region bounded by conducting walls that supports standing electromagnetic waves at specific frequencies. When you drive a cavity at its resonant frequency, energy builds up inside. The driving source keeps pumping energy in, and losses (from wall resistance, coupling ports, and field absorption) take energy out. At steady state, the energy stored in the cavity is much larger than the energy flowing through it at any instant.

How much larger? That's the Q factor. A cavity with $Q = 10,000$ stores 10,000 times the energy that passes through it per radian of oscillation. At a microwave frequency of 10 GHz with 1 watt of input power and $Q = 10,000$, the internal stored energy corresponds to a field intensity roughly 10,000 times what you'd get from that 1 watt propagating freely. The radiation pressure on the cavity walls scales with the square of the field amplitude, so it goes up by a factor proportional to Q .

The walls don't just contain the field. They reflect momentum back into the interior. Every reflection is a momentum reversal, which means the wall absorbs a double kick, once from the incoming wave, once from the outgoing wave. In a high-Q cavity, those kicks add up fast.

Quality Factor as a Compression Metric

Q factor is defined as 2π times the ratio of energy stored to energy dissipated per cycle. For a microwave cavity, Q values of 10^3 to 10^6 are achievable depending on material and geometry. Superconducting cavities used in particle accelerators reach Q values above 10^{10} . At those levels, the field compression inside the cavity is extraordinary.

From a momentum engineering standpoint, Q is your primary design parameter for force generation. Everything else being equal, a factor-of-ten increase in Q gives you a factor-of-ten increase in internal field energy, and a factor-of-ten increase in radiation pressure on the cavity walls and any contents.

This is why material choice matters so much for EM momentum engineering. Copper cavities at room temperature get you into the thousands of Q. Electropolished copper or silver gets higher. Superconducting niobium gets you to millions or more. The tradeoff is complexity and cost. For near-term applications, room-temperature high-Q cavities are practical and sufficient for micronewton-to-millinewton force levels.

Phased Array Superposition

A single field source at power P produces field amplitude proportional to $\sqrt{P}$. If you add N coherent sources in phase, the total field amplitude is N times a single source's amplitude. Energy density goes as amplitude squared, so it scales as N^2 . Force scales as energy density, so it also scales as N^2 .

With N independent sources at total power $N \times P$, a simple incoherent arrangement gives you N times the force. A coherent phased arrangement gives you N^2 times the force. That's the difference between illumination and compression. Phase coherence turns N sources into a coherent pump that concentrates momentum into the target volume rather than spreading it isotropically.

This is not a new principle. It's the foundation of phased array radar, directed energy weapons, and high-intensity focused ultrasound. The EM momentum engineering angle is to use that coherent concentration not for signal transmission or heating, but specifically for force production at the focal point.

The EmDrive as a Case Study

The EmDrive controversy is instructive, whatever you think of the specific claims. The device is a tapered resonant cavity, frustum-shaped, with a larger face and a smaller face, driven at microwave frequencies. The claim was that it produced a small net thrust without expelling propellant.

The controversy exists because momentum conservation in a closed resonant system is genuinely subtle. Standard analysis says you can't get net thrust from a closed cavity because the radiation pressure on the large end and the small end cancel when you integrate over the whole system. But the geometry of a tapered cavity means the field distribution inside is not symmetric, field intensity is higher at the small end because the cavity is more confining there. The radiation pressure on the small end wall is therefore higher than on the large end wall.

Whether this produces net thrust depends critically on what happens to the momentum of the field at each reflection, how the electromagnetic momentum in the cavity couples to mechanical momentum of the cavity structure, and what role the surrounding medium (vacuum or otherwise) plays in the accounting. These are not trivial questions. The reason to look at the EmDrive is not as a solved engineering problem but as an example of a genuinely hard momentum accounting problem in a resonant system. The right response is better measurement, not dismissal.

Compression Geometry: Cylindrical vs. Conical

Cavity geometry determines mode structure, which determines where the field is most intense, which determines where the force is largest. A cylindrical cavity has uniform cross-section and supports modes with relatively uniform field distribution. A conical or tapered cavity breaks that symmetry. The field concentrates toward the narrow end as the geometry tightens.

For force production, you want the maximum energy density at the face you're trying to push or the surface you're trying to load. Conical geometries, frustum shapes, and tapered structures are natural momentum concentrators. Horn antennas, tapering waveguides, and conical resonators are all variants on this principle.

CHAPTER FIVE

Material Steering: Making the Field Go Where You Want

Conductors Are Not the Only Tool

Most engineers think of EM field control in terms of conductors: metal walls, ground planes, waveguide boundaries. Conductors are powerful because they enforce a hard boundary, the tangential electric field goes to zero at a perfect conductor surface. But they're not the only field-shaping tool, and for momentum engineering purposes, they're often not the right one.

The reason is that a conductor reflects the field. Reflection is momentum transfer, which is useful, but a conductor doesn't let you shape the internal distribution of the field without changing the physical geometry of the conductor. Dielectrics and metamaterials let you steer the field inside a volume, changing where the momentum lives, without moving the boundaries.

Dielectric Lensing for Momentum Purposes

A dielectric material has a permittivity higher than free space. When a field enters a high-permittivity region, it slows down. The wavelength shortens. The field bends toward the normal at the interface (Snell's law for EM waves). For a curved dielectric surface, this produces lensing, convergent or divergent wavefronts.

Dielectric lenses are everywhere in RF and microwave engineering. They're used in antenna feedhorns, millimeter-wave imaging systems, and satellite communication terminals. The design objective is usually beam shaping for signal transmission. The same geometry can be used with a different objective: steer the wavefront so that the maximum momentum flux is directed at a target.

Specifically, if you design a dielectric lens to create a converging wavefront that focuses at a small spot, you've created a region of high momentum density at that spot. The field's momentum, which was distributed over the lens aperture, is now concentrated at the focus. Any object placed at the focus experiences the combined radiation pressure of all the field that converged there. This is the macro-scale equivalent of an optical tweezer.

Metamaterials: Engineering the Effective Medium

Metamaterials are structures with sub-wavelength feature sizes that behave, to the EM field, like a homogeneous medium with properties not found in natural materials. The effective permittivity and permeability of the structure are determined by the geometry of the unit cells, not by the atomic properties of the constituent materials.

This gives you access to parameter space that natural materials can't reach. Negative effective permittivity (common in metals above the plasma frequency). Negative effective permeability (achievable with resonant structures like split-ring resonators). Negative index of refraction, where both ϵ and μ are simultaneously negative, the field bends the wrong way at an interface, in the sense of backward refraction. Zero-index materials, where the phase velocity goes to infinity and the field propagates with no spatial phase variation.

These aren't science fiction. Negative index metamaterials were demonstrated experimentally in the early 2000s and have been commercially produced in the microwave range. Zero-index metamaterials have

been demonstrated in the optical regime. The fabrication challenge is that the feature size has to be smaller than the wavelength, so as you scale down in wavelength (up in frequency), the feature size requirement gets harder to meet. At microwave frequencies, the features are millimeter-scale, which is straightforward with standard PCB and machining technology.

What Metamaterials Let You Do for Momentum Engineering

Negative index materials redirect field momentum in counterintuitive directions. A field entering a negative-index slab at normal incidence doesn't just slow down, it effectively propagates backward in the sense that the phase front moves in the direction of energy flow, but the wavevector points the other way. This lets you design momentum-redirecting structures that would be physically impossible with natural materials.

Zero-index materials are interesting for a different reason. When the index is zero, the field sees no spatial variation, it's as if the interior of the material is a single point in phase space. Field configurations that would normally produce internal reflections and standing waves instead pass through without standing-wave structure. This can be used to reduce parasitic internal forces and direct net momentum toward a specific output face.

Anisotropic metamaterials, with different effective index along different axes, can route field energy along preferred paths inside a structure, the same way an optical fiber routes light along its axis. But unlike a fiber, a metamaterial structure can route energy in a volumetric pattern, not just a linear one.

The Momentum Funnel: Substrate-Layer Engineering

The most powerful approach combines dielectric and metamaterial steering with cavity geometry to create a momentum funnel: a structure where field energy enters from a broad face, is directed and compressed through a designed material stack, and exits (or concentrates) at a narrow target face with high momentum density.

The key concept is impedance transformation through layered media. A stack of alternating high-permittivity and low-permittivity layers acts as a spatial transformer. Each layer interface is a partial reflector. When the layers are designed with the right thicknesses and index values, constructive

interference concentrates field energy at one end of the stack, similar to how a quarter-wave transformer in transmission line theory matches impedances and maximizes power transfer.

The same principle works in acoustics, acoustic impedance matching stacks are used in ultrasound transducers to couple energy from piezoelectric elements (high acoustic impedance) into tissue (low acoustic impedance). The EM version works the same way, and the design equations are structurally identical.

The Steering Equation in Plain Terms

The net force on a volume element of a dielectric in an EM field is given by the gradient of field energy density multiplied by the susceptibility of the material. If you have a spatial gradient of field energy density, higher on one side of the object than the other, the object feels a net force toward the high-field region. Make that gradient steep and pointed, and you have a force that's real and measurable.

The engineering objective: design a material structure and field geometry such that the energy density gradient at the target location is as steep as possible and points in the direction you want. Everything else follows from the Maxwell stress tensor.

CHAPTER SIX

Field Geometry as a Force Generator

Putting the Pieces Together

Chapters 3 through 5 laid out three tools: lattice geometry (the structure of the propagation environment), EM compression (concentrating field energy), and material steering (directing where that energy lives). This chapter is about how you combine them into a system that produces a sustained, directional mechanical force.

The objective is simple to state. You want a configuration where the Maxwell stress tensor is asymmetric, where the electromagnetic stress on one face of a system is systematically higher than on the opposing

face. Integrate that difference over the system volume, and you get net force. Keep the field running, and you get sustained thrust.

The Maxwell Stress Tensor as a Design Target

The Maxwell stress tensor T is a 3×3 matrix where component T_{ij} gives the flux of the i -th component of momentum in the j -th direction. The diagonal components (T_{xx} , T_{yy} , T_{zz}) represent pressure, force per unit area pushing perpendicular to a surface. The off-diagonal components represent shear stresses.

For a simple plane wave propagating in the z -direction, the only nonzero stress is T_{zz} , which equals the energy density u . The wave pushes in the direction it's moving with a pressure equal to its energy density. This is the radiation pressure formula: $P_{\text{rad}} = u$ for an absorber, $2u$ for a reflector.

For a more complex field configuration, a standing wave, a mode in a cavity, a superposition of multiple waves, all components of the stress tensor can be nonzero and spatially varying. The net force on an object is the integral of the tensor over a closed surface surrounding the object. Asymmetry in the field configuration produces asymmetry in the integral produces net force.

Design Approach: Asymmetric Resonators

The most straightforward force-producing geometry is an asymmetric resonator: a cavity where the field distribution is intentionally non-uniform, with higher energy density at one end than the other. The tapered cavity described in Chapter 4 is one version. But you can achieve the same effect through other means.

Off-center field concentration: place a high-permittivity insert inside a symmetric cavity, positioned toward one end. The insert shifts the field distribution toward itself. The end of the cavity nearest the insert now has higher field intensity than the opposite end. Radiation pressure on the near end is higher. Net force acts on the insert and, through the insert, on the cavity as a whole.

Asymmetric mode coupling: a cavity can be designed to preferentially excite modes with a known field asymmetry. By choosing the coupling port location, the mode frequency, and the cavity geometry together, you select a mode where energy is concentrated in a specific region rather than distributed uniformly.

Impedance Mismatches as Force Sites

Every time an EM field crosses a material boundary, conductor surface, dielectric interface, cavity end wall, there's an impedance mismatch. Some of the field reflects. Some transmits. The momentum accounting at that boundary produces a force on the interface.

Sharper impedance mismatches mean more reflection, which means more momentum reversal, which means more force. If you design a system where all the sharp impedance mismatches are on one side of a structure, all the force kicks are on that side. The other side has smooth, gradual impedance transitions with minimal momentum transfer. You've created a force asymmetry through deliberate impedance engineering.

This is exactly what happens at the small end of a tapered cavity (or a conical horn). The field hits the end wall at high intensity with a sharp impedance step (air to metal conductor), and all that momentum is transferred to the wall. The field at the large end sees a gradual taper with lower field intensity. The force budget is not balanced, and the imbalance is by design.

Worked Example: Tapered Dielectric Rod in a Resonant Cavity

Place a dielectric rod inside a resonant cylindrical cavity. The rod runs along the axis of the cavity but is tapered, it has higher permittivity (or larger cross-section) at one end and lower permittivity at the other. The cavity is driven at a frequency where the mode field is strongest along the axis.

The high-permittivity end of the rod attracts field energy. The field is more intense at that end. The gradient dielectric force on the rod (from the gradient of E^2 times the permittivity variation) points from the sparse end toward the dense end. There's also a direct radiation pressure on each end face of the rod, higher at the dense end because field intensity is higher there.

The rod experiences a net axial force in the direction from sparse to dense end. This is not a thought experiment. The physics is standard. The force is calculable from the Maxwell stress tensor. The magnitude depends on Q, drive power, rod geometry, and permittivity contrast.

Scaling and Force Magnitudes

Radiation pressure scales as field energy density divided by c . Field energy density in a resonant cavity scales as Q times input power divided by cavity volume. So net force scales roughly as:

$$F \approx (Q \times P_{\text{in}}) / (V_{\text{cavity}} \times c)$$

At $Q = 10,000$, input power of 1 kW, cavity volume of 1 liter (10^{-3} m^3), force is approximately $10,000 \times 1,000 / (10^{-3} \times 3 \times 10^8) \approx 33$ micronewtons. At $Q = 100,000$ and 1 kW input, it's 330 micronewtons. At 10 kW, it's 3.3 millinewtons.

These aren't large forces. A paperclip weighs about 4 millinewtons. But millinewton-scale forces are precisely the range needed for spacecraft station-keeping, attitude control, and precision positioning of small objects in a lab or manufacturing environment. The scale is not the limitation. The scale is the application space.

CHAPTER SEVEN

Practical Engineering Applications

The previous six chapters built the conceptual and theoretical foundation. This chapter is the payoff. Here are the real and near-real applications of EM momentum engineering, laid out by domain, without hype and without excessive qualification.

Spacecraft and Orbit

The most compelling near-term application for EM momentum engineering is spacecraft propulsion and attitude control, specifically in regimes where mass is the limiting resource and power is relatively available.

Station-keeping for a geostationary satellite requires millinewton-scale thrust for months to years. Current ion thrusters and Hall effect thrusters do this job well, but they require propellant, xenon, krypton, argon, and propellant is mass that has to be launched. Every kilogram of propellant launched to GEO costs roughly \$20,000 to \$30,000 in launch costs (and more to high-energy orbits). An EM momentum thruster

that runs on electrical power has zero propellant budget, only an electrical power budget, which on a solar-powered satellite is not the limiting constraint.

Solar radiation pressure is already used passively. Solar sails use it for primary propulsion on small spacecraft. Active EM pressure, generating the field on-board rather than receiving sunlight, allows the force direction and magnitude to be controlled independently of the spacecraft's orientation relative to the sun. That's operationally valuable for attitude control and precise orbit maintenance.

The force levels from an on-board EM momentum device at realistic power and Q levels are in the micronewton to millinewton range, compatible with attitude control and drag compensation in low Earth orbit, where atmospheric drag produces deceleration forces in the same range at altitudes above 400 km. No propellant, no exhaust plume (which contaminates sensors), no chemical hazard, solid-state operation.

Precision Positioning

Optical tweezers use focused laser light to trap and move microscopic particles, cells, organelles, colloidal beads, single molecules. The force is the gradient force from the radiation pressure of a tightly focused Gaussian beam. The technology is well-established; Arthur Ashkin won the Nobel Prize in Physics for it in 2018.

The same principle scales to larger objects at lower frequencies. Shaped microwave or millimeter-wave fields can exert controlled forces on objects in the gram-to-kilogram range, depending on power level and geometry. This is the macro-scale optical tweezer: a contactless positioning system that uses field geometry, not mechanical contact, to move and hold objects.

Applications in semiconductor manufacturing, cleanroom handling of sensitive components, and biological sample preparation are all technically feasible. The key advantage is zero contact, the field touches nothing, leaves no residue, doesn't scratch surfaces, doesn't transfer contaminants. In environments where contact contamination is a hard limit, a field-based manipulator is superior to any mechanical gripper regardless of its force resolution.

Field stability is the engineering challenge. To position an object to nanometer precision, the field geometry has to be stable to a corresponding precision. That requires temperature-controlled cavities,

stable power supplies, and vibration isolation. All of these are standard metrology problems with known solutions.

Material Processing

High-intensity EM fields do more than heat things. They interact with material microstructure directly through the mechanisms described in earlier chapters: gradient forces on dielectric inclusions, radiation pressure on grain boundaries, anisotropic field-induced stress that can bias crystallographic alignment.

Field-assisted sintering is already industrial. The FAST (Field Assisted Sintering Technology) process, also called spark plasma sintering, uses applied fields during powder consolidation to produce denser, more homogeneous ceramic and metal parts at lower temperatures and shorter times than conventional sintering. The mechanism involves both Joule heating and direct field-matter interaction at grain boundaries.

Extending this concept: a high-intensity focused EM field applied to a crystalline material during processing can bias grain growth direction, influence residual stress state, and modify defect density. This isn't theoretical, there's a body of experimental literature on field-assisted phase transformation and texture development in metals and ceramics. The momentum engineering angle adds directionality to that interaction: instead of just applying a uniform field, you apply a field with a designed momentum gradient, so the field-matter interaction has a preferred direction and a quantifiable force component.

Residual stress relief is another target. Vibration stress relief (VSR) uses mechanical vibration to relax residual stresses in weldments and castings. EM stress relief, using directed field momentum to substitute for or augment mechanical vibration, is a logical extension. The advantage is that the "vibration" can be at frequencies and with spatial patterns that mechanical exciters can't achieve.

Shielding and Wavefront Sculpting

Conventional EM shielding absorbs or reflects incoming energy. Absorbing shielding turns field momentum into heat in the absorber, which creates a thermal load. Reflecting shielding sends the momentum back, which creates a radar return and doesn't protect the shielded object from the direction it came from.

Active shielding through EM momentum engineering does something different: it redirects incoming field energy. Instead of absorbing or reflecting it, the shield deflects the wavefront around the protected region, so the field energy continues in a new direction that doesn't intersect the target. This is the EM equivalent of a fluid dynamic fairing, the field flows around the obstacle rather than being stopped by it.

Phase-conjugate mirrors are a specific tool for this. A phase-conjugate mirror is a structure that, when illuminated by a wave, returns a wave that is the time-reversed copy of the incident wave. The reflected wave retraces the exact path of the incoming wave back to its source, carrying the momentum back the way it came. This effectively returns incoming radar energy to its source with the same direction it arrived from, a powerful countermeasure.

Phase conjugation has been demonstrated experimentally in multiple regimes, optical, acoustic, and microwave, using nonlinear media and specifically designed resonant structures. Engineering the phase-conjugate response into a practical panel or antenna system is a current research frontier, not a solved commercial product, but the underlying physics is established.

Power Transfer Without Wires

Wireless power transfer is already commercial. Inductive charging (the Qi standard) transfers power at short range through resonant magnetic coupling. It works at distances of a few centimeters with efficiencies above 90%. It's in every modern smartphone.

EM momentum engineering extends this concept in two directions: longer range and higher directionality. Standard resonant inductive coupling is near-field, the field doesn't propagate; it stores energy in a reactive near-field that couples between source and receiver at close range. For distances beyond a few wavelengths, you need to switch to a propagating wave approach: a directed EM beam aimed at a rectifying antenna (rectenna) on the receiving end.

The rectenna concept is old. William Brown demonstrated microwave power beaming with a rectenna in the 1960s. Modern high-efficiency rectenna arrays at microwave frequencies achieve DC conversion efficiencies above 80%. The beam delivery problem, getting the power to the rectenna with minimum spillage and losses, is where EM momentum engineering is relevant. A phased array transmitter designed to produce a coherent, focused momentum beam at the rectenna location delivers more power to the target per watt of transmitter power than a simple antenna does.

Applications include wireless power delivery to remote sensors (eliminating battery replacement), power beaming to small spacecraft from a ground station or orbital relay, and charging of autonomous vehicles or robotic systems without physical docking. None of these require new physics. They require engineered field geometry and good rectenna design.

CHAPTER EIGHT

Measurement and Verification: How You Know It's Working

The Measurement Problem

EM force measurements are hard. The forces are small, micronewtons to millinewtons, and the physical environment around any EM source is full of confounders: Joule heating that causes thermal expansion, Lorentz forces on current-carrying elements, vibration from power supplies and fans, and ambient acoustic noise. All of these can produce mechanical signals that look like thrust if you're not careful.

The history of propulsion experiments is littered with false positives caused by inadequate systematic error control. An experiment that claims millinewton-level thrust from an EM device is not interesting until it has ruled out all the mundane explanations. This isn't skepticism for its own sake, it's the standard of evidence required to make the result useful. Here's how you do it right.

The Torsion Balance

The torsion balance is the gold standard for small force measurement. Suspend the test device from a thin fiber, fused silica or tungsten, so it can rotate about the vertical axis. An applied horizontal force produces a torque, which twists the fiber, which you measure as an angular displacement. The restoring torque of the fiber is known from calibration, so displacement maps directly to force.

This is the same technique Cavendish used to measure the gravitational constant in 1798. It's been refined considerably since then. Modern torsion balances in metrology labs can resolve forces below 1 nanoNewton. The technique is limited by fiber thermal noise (Brownian motion of the fiber), acoustic

noise coupling, and residual air currents. In a well-designed setup, resolution to tens of nanonewtons is achievable at room temperature.

For a device producing micronewton-scale forces, a torsion balance with 10 nanonewton resolution gives you a signal-to-noise ratio of at least 100. That's enough to see the signal clearly, measure its dependence on power level, and perform null tests.

Laser Interferometry

If the torsion balance displacement is too small to read optically, a laser interferometer on the fiber head solves the problem. Modern Michelson interferometers with frequency-stabilized lasers resolve displacements to picometer scales. Even for extremely small radiation pressure forces, the resulting displacement of a torsion balance is detectable.

LIGO (the Laser Interferometer Gravitational-Wave Observatory) uses this principle to detect displacements smaller than 10^{-19} meters. You don't need LIGO-level sensitivity for EM momentum experiments, but the same fundamental approach, laser interferometry on a compliant mechanical element, is the right tool for detecting small systematic displacements.

Systematic Error Control

Thermal effects are the most common false positive source. Asymmetric heating of the test device causes thermal expansion, which moves the device. If the EM source heats the device asymmetrically, it produces a displacement that correlates with RF power level and looks exactly like thrust. The countermeasure is careful thermal mapping: measure the temperature distribution on the device under test at every power level, calculate the expected thermal deformation, and subtract it from the observed displacement.

Lorentz forces on current-carrying elements are the second common confounder. Any cable carrying RF power to the test device in an ambient magnetic field experiences a force. The fix is to run the cables symmetrically and to cancel the DC magnetic environment, or to use flex cables with negligible mechanical stiffness.

Acoustic coupling from the power source is the third. Any mechanical vibration in the power supply that reaches the test device through the support structure will show up as noise. Acoustic isolation stages, vibration-damped support pads, bellows-isolated cables, isolation tables, are standard solutions.

Null Tests

The null test is the most powerful tool in the measurement toolkit. Build a symmetric version of the test device: same geometry, same materials, same internal field distribution, but with the asymmetry deliberately removed. If the symmetric device still shows a force signal, the signal is a systematic error, not physics. If it goes to zero (within noise), the systematic errors are under control, and the asymmetric device's signal is real.

For a tapered cavity experiment, the null device is a cylindrical cavity (uniform cross-section) driven at the same power level. It has no geometric momentum asymmetry and should produce zero net force. If it does, great, you've shown your measurement system is clean. If it shows force, something is wrong and you need to find it before claiming any results from the tapered version.

A second type of null test is the off-resonance test. Drive the device at a frequency far from resonance, where Q is low and field intensity is minimal. The systematic errors (thermal expansion, Lorentz forces) scale with input power, not with field intensity. The EM momentum forces scale with stored field energy, which is Q times higher at resonance. So at off-resonance, you should see systematic errors but no physics signal. At resonance, you should see the physics signal above the systematic baseline. If the resonance condition produces additional force that scales correctly with Q , that's your EM momentum signal.

The Measurement Signature

What you're looking for is a correlation between stored field energy (not just input power) and mechanical force, that survives all null tests and all off-resonance baselines. The signal should scale with $Q \times P_{in}$ for a resonant device. It should reverse direction when you reverse the device orientation. It should disappear when you eliminate the geometric asymmetry. It should be frequency-selective, peaking at cavity resonance.

That combination of signatures is unambiguous. You're not guessing about systematic errors, you're testing for a specific predicted pattern and showing that the pattern matches. That's the measurement standard the field needs to move forward.

CHAPTER NINE

Where This Goes: Open Problems and Engineering Frontiers

An Honest Assessment

Not everything in this book is equally settled. Some of it is textbook physics. Some of it is engineering extrapolation from well-established principles. Some of it is speculative in the specific sense that the theoretical framework is motivated and internally consistent, but the experimental confirmation is incomplete. Knowing which is which is essential for doing productive work in this area.

What Is Settled

Radiation pressure is real. It's been measured to high precision since 1903, and every modern experiment that's looked for it has found it. You cannot build an argument against radiation pressure, it would require rewriting Maxwell's equations and the special theory of relativity.

Resonant cavities amplify field energy density by a factor proportional to Q . This is engineering fact. Every microwave cavity designer knows it. Every accelerator physicist depends on it. The Q values achieved by superconducting cavities, above 10^{10} , are a matter of routine measurement.

Metamaterials steer fields in ways natural materials cannot. Negative index refraction has been measured. Zero-index propagation has been demonstrated. Transformation optics devices (EM cloaks, field concentrators) have been built and tested at microwave frequencies. The fabrication physics is understood and controllable.

EM momentum transfer produces measurable forces on surfaces, materials, and particles. Optical tweezers are a commercial product used in biology labs worldwide. Laser cooling of atoms, using radiation pressure to slow atomic motion to microkelvins, is the foundation of atomic clocks and quantum computing platforms. These aren't marginal demonstrations. They're precision instruments.

What Is Open

Net thrust from a closed resonant cavity remains contested. The theoretical analysis is not fully settled. The momentum accounting in a system that includes the electromagnetic field, the cavity structure, the feed lines, the surrounding medium, and any evanescent fields extending outside the cavity involves subtleties that honest theorists disagree about. Some analyses find net thrust as an artifact of Abraham-Minkowski momentum formulation choices. Others find it as a real effect. Others find no effect at all and attribute all experimental positive results to systematic errors.

The experimental record is equally murky. Several groups have reported positive results from tapered cavities. Several others have reported null results with what they consider better systematic error control. The data are not yet clean enough to settle the question, which is frustrating but also the normal state of a frontier measurement science.

The right response to this ambiguity is not to pick a side and defend it. The right response is to build better experiments.

The UST Frontier: Substrate-Coupled Force

The speculative frontier, motivating from the UST framework, is the question of substrate-coupled momentum transfer. In standard physics, when an EM field exerts a force on a surface, the momentum balance is satisfied by photon recoil: the field carries momentum in one direction, the surface receives momentum in the other direction, and the total momentum of the system (field plus matter) is conserved.

In the substrate picture, there's an additional possibility: the field couples to the substrate geometry, and the reaction force propagates through the substrate rather than through photon recoil. If this is correct, the momentum accounting for a closed cavity changes. The cavity doesn't have to emit radiation to conserve momentum, it exchanges momentum with the substrate directly, and the substrate is (for practical purposes) an infinite reservoir.

This is a strong claim that requires strong evidence. The theoretical motivation is that the substrate is a physical medium with defined mechanical properties, and coupling between the field and the medium is not just conceptually possible but physically expected given that ϵ_0 and μ_0 are properties of the medium. Whether the coupling is strong enough to produce engineering-relevant forces at realistic power levels is not theoretically determined without a more specific model of the substrate's mechanical response.

The experimental test for substrate coupling is, in principle, the same torsion balance measurement described in Chapter 8. The signature of substrate coupling would be a force that doesn't scale with photon momentum flux (doesn't follow P/c) but scales differently with frequency, geometry, or material composition. If you see a force that can't be accounted for by standard radiation pressure, and it survives all null tests, and it correlates with something the substrate coupling picture predicts, that's your signal. Building that experiment is the near-term research agenda.

The Engineering Agenda in Order of Priority

First priority: measurement infrastructure. High-Q room-temperature resonators. Torsion balance test stands with interferometric readout. Well-shielded test environments with documented thermal and acoustic noise floors. Without this infrastructure, no theoretical debate can be resolved. Build the instruments first.

Second priority: demonstrate EM momentum force generation at the millinewton level with verified null tests. This is conservative engineering using established physics. It doesn't require the UST extension or substrate coupling. It just requires a well-designed asymmetric resonant structure and careful measurement. This is achievable with existing technology and modest resources.

Third priority: develop the application-specific engineering for the near-term markets, precision positioning, drag compensation, contactless manipulation. Get something useful in the field. Build the engineering base that supports more ambitious research.

Fourth priority: when the measurement infrastructure is mature, design the substrate-coupling test explicitly. Control for all standard momentum transfer mechanisms, look for a residual, characterize it, and publish the data regardless of outcome.

Summary: The Engineering Mindset for EM Momentum

What You Now Know

EM fields carry momentum. The Poynting vector gives you the momentum flux at every point in the field. The Maxwell stress tensor gives you the force per unit area on any surface due to the field. These are exact, classical results that have been verified experimentally for over a century. There is no uncertainty here. The momentum is real, and it pushes things.

The substrate perspective, treating the vacuum as a structured medium rather than a featureless backdrop, gives you a richer engineering vocabulary. Local field geometry affects propagation, field intensity, and momentum distribution. Shaping that geometry means shaping the force output of your system. Whether you accept the full UST framework or work purely within standard electrodynamics, the design principle is the same: engineer the environment the field propagates through, and you engineer the force the field produces.

The Four Engineering Tools

Everything in EM momentum engineering comes down to four handles. Get these right, and the physics takes care of itself.

Geometry. The shape of the cavity, the waveguide, the antenna, and the material stack determines where the field is intense, where the momentum gradients are steep, and in which direction the net force points. Geometry is your primary tool and your most powerful one. You don't need exotic materials to produce asymmetric stress distributions, you need the right shape.

Materials. Permittivity and permeability determine how the field interacts with matter. High-permittivity materials attract field energy and concentrate it. Metamaterials extend the accessible parameter space to regions natural materials can't reach. Impedance mismatches are force sites, engineering where those mismatches occur means engineering where the momentum transfers.

Power. Field energy density determines the magnitude of radiation pressure. More input power means higher field energy density (in proportion to Q), which means higher force. Power is the throttle, not the steering. Design the geometry and materials first, then scale with power.

Coherence. Phase alignment between multiple field sources multiplies force nonlinearly. N coherent sources produce N^2 times the force of a single source at the same total power. Coherence is the difference between illumination and pressure. Phased arrays, resonant buildup, and mode selection are all coherence tools.

The Discipline in Context

EM momentum engineering is not a new branch of physics. It's an extension of antenna engineering, waveguide design, and optical engineering into the force domain. The mathematical tools are the same. The measurement techniques are the same. The material physics is the same. What's different is the objective function: instead of optimizing for radiated power, signal-to-noise ratio, or image resolution, you're optimizing for mechanical force.

That shift in objective function changes which design choices matter and which don't. A resonant cavity with high Q and asymmetric field distribution is a good antenna in some configurations and a terrible one in others, but it's a good force generator regardless. The engineering skill is in knowing which parameters to optimize for which objective, and building the right objective into the design from the start.

Why This Field Is Underdeveloped

EM momentum engineering sits at the intersection of plasma physics, microwave engineering, and precision mechanics. The plasma physics community studies field-matter interaction but isn't focused on force generation from resonant cavities. The microwave engineering community optimizes for signal transmission and reception, not mechanical force output. The precision mechanics community measures small forces with extraordinary skill but usually in the absence of strong EM fields.

Nobody owns this space. There's no professional society for it, no dedicated journal, no established graduate curriculum. The applications span aerospace, manufacturing, metrology, and energy, which means no single industry has had the incentive to develop it systematically.

That's the opportunity. The foundational physics is established. The tools exist. The application domains are real and commercially significant. The field is waiting for engineers who are willing to operate at the intersection of disciplines, use the established tools from all three communities, and build the measurement infrastructure that turns theoretical motivation into engineering reality.

The Call to Action

Start with what's unambiguously real: radiation pressure from a high-Q resonant cavity is calculable, measurable, and achievable with room-temperature hardware at kilowatt power levels. Build that experiment. Measure the force. Do the null tests. Publish the data. That's a contribution even before you get to the more speculative questions about substrate coupling and closed-cavity propulsion.

Then extend. Design an asymmetric resonator. Optimize the geometry for force rather than signal. Test it against the symmetric null. If the signal survives, you're in business. If it doesn't, you've learned something about systematic errors and you're better positioned to design the next version.

The physical tools, high-Q cavities, laser interferometers, torsion balances, metamaterial fabrication, are all mature technologies. The missing piece is focused engineering attention. Point that attention at the right problem, and the progress will be fast. Fields carry momentum. Momentum can be transferred. Transfer can be engineered. The rest is details.