Cosmological Bounce and Emergent Particle Structure

By Dustin Lee

Introduction

All known stable particles do not emerge from quantum field fluctuations in a vacuum, but from the geometric deformation of a pre-existing substrate during and immediately after the cosmological bounce. The substrate is a continuous, compressible, topologically rich medium that carries the electromagnetic field as its primary stress mode. The standard model vacuum is not nothing; it is a geometric medium whose structure during the bounce determines everything that follows. Prime-number subdivisions of geometric knot structures naturally select stable particle configurations, providing a first-principles derivation of why certain particles exist, why they carry the masses and charges they do, and why all others decay. Container geometry, the large-scale structure of space, emerges from the same substrate dynamics that produce particles. The Cosmic Microwave Background is reinterpreted as the fossil record of the first substrate oscillation wave: a direct geometric imprint of the bounce, not a thermal relic in the standard sense. This paper lays out the complete causal sequence: bounce → EM field → substrate modes → photons → electrons → protons → neutrons → container geometry → CMB. The framework requires three substrate parameters and derives from them the whole of particle physics and cosmology, with no additional free parameters and several distinguishing observational predictions.

The Problem with Standard Particle Creation

The standard model of particle physics is the most precisely tested theoretical framework in the history of science. I do not dispute its predictions. What I dispute is its explanatory depth. The standard model describes particles with extraordinary accuracy; it does not explain why those particular particles exist, why they carry the specific masses and charges they do, or what mechanism generated them at the origin of the universe. These are not pedantic questions. They are the central questions of fundamental physics, and the standard account defers them without resolution.

Particles in the standard cosmological picture are created from quantum field fluctuations in the vacuum. The vacuum, in this account, is a seething quantum medium whose fluctuations are promoted to real particles by the dynamics of inflation. But this immediately raises a question that the framework simply does not answer: what is the vacuum? The word is used as though it designates something geometrically featureless, an absence. In reality, the quantum vacuum is already assumed to be a structured medium, one that supports fields, that has a defined energy density, that carries representations of the Lorentz group, and that undergoes fluctuations with a specific spectral character. This is not nothing. It is a highly structured physical substrate. The standard treatment names it, uses it, and then declines to say what it is made of or what its geometry is.

Inflation is the dominant mechanism invoked to explain the initial conditions of the universe: the flatness, the horizon, and the large-scale homogeneity. I want to be direct about inflation's status. Inflation is not a mechanism in the same sense that the bounce is a mechanism. Inflation is a mathematical degree of freedom, a scalar field with a suitably flat potential, added to general relativity in order to solve specific observational puzzles. It does not derive the masses of particles. It does not derive electric charge. It does not explain spin. It assumes all of these as pre-given and uses them. The particle content of the universe after inflation is identical to the particle content assumed before inflation started. Inflation explains the geometry of initial conditions without explaining the content.

My approach is not to reject quantum mechanics. Quantum mechanics is correct. What I am doing is going deeper: asking what the physical substrate of quantum fields actually is, and what its geometry does during and after the cosmological bounce. The quantum formalism describes the behavior of the substrate modes correctly. The substrate framework explains what those modes are.

The bounce is real. Loop quantum cosmology, bouncing cosmology, and a range of pre-big-bang models have established with increasing mathematical rigor that the classical Big Bang singularity is not a physical event, it is an artifact of applying general relativity beyond its domain of validity. The physical event is a bounce: a maximum-density, minimum-volume state from which the universe rebounds. I accept the bounce as the unambiguous starting condition. The question is not whether the bounce occurred. The question is what the bounce does to the substrate. What is the geometric structure of the quantum vacuum substrate, and what does the cosmological bounce do to it? The answer to that question is this entire paper.

The Substrate

The substrate is the physical medium that underlies all fields and particles. It is continuous, compressible, and capable of sustaining topological deformation. It is not a collection of particles, it has no discrete constituents. It is a genuine continuum: a medium with curvature, tension, and a characteristic compression modulus that determines how it resists compression and how it rebounds when compressed.

Before the bounce, the substrate contains no particles. It contains geometry: it has a large-scale curvature (its container shape), local tension (its resistance to shear), and topological features carried over from the prior contraction cycle. These prior-cycle features are not erased by the bounce. The substrate has topological memory, geometric structures that persist through the maximum-compression event. This is crucial, and I will return to it in the discussion of the CMB.

Let me be equally explicit about what the substrate is not. It is not aether. The nineteenth-century luminiferous aether was conceived as a static, passive background medium whose only function was to carry light waves. The substrate is nothing like that. It is dynamic: it stretches, folds, knots, twists, and oscillates. It is not a background against which events occur. It is the events. The substrate does not sit still while fields propagate through it. The substrate's deformation modes are the fields. There is no separate ontological layer. There is only the substrate and its geometry.

The electromagnetic field is the substrate's primary stress mode. When the substrate is sheared or compressed in a transverse direction, what we measure as an electric or magnetic field is the substrate responding to that deformation. The electromagnetic field is not a thing living in space. It is space (the substrate) under a specific class of stress. This identification is not metaphorical. The mathematical structure of the substrate's stress tensor, under transverse deformation, is precisely the structure of Maxwell's equations. The EM field does not live in the substrate; it is the substrate in shear.

Gravity operates differently. Gravity is not a substrate stress mode in the same transverse sense as electromagnetism. Gravity is the large-scale curvature shape of the substrate container. The macroscopic geometry of the substrate as a whole. I will develop this distinction carefully in Section 10. The key point here is that electromagnetism and gravity, while both geometric in origin, correspond to different scales and modes of substrate deformation: EM is local shear stress, gravity is global container curvature.

The substrate has a minimum compression length. This is not an arbitrary ultraviolet cutoff imposed by hand to regulate divergences. It is a geometric consequence of how a continuous medium with finite tension resists compression. A medium with finite compression modulus cannot be compressed below a characteristic length determined by the ratio of its tension to its modulus. This minimum length is the Planck length. The Planck scale is not a fundamental input; it is a derived property of the substrate's mechanical parameters. Below the Planck scale, the substrate cannot be further compressed; the concept of a smooth, continuously deformable medium breaks down. This is why quantum gravity diverges at the Planck scale: the continuous substrate description stops being valid there.

The Bounce

At the bounce, the entire substrate reaches its state of maximum compression. This is not a singularity. It is not a mathematical breakdown. It is a physical maximum-density state, constrained by the substrate's compression modulus. The substrate cannot be compressed further because its elastic resistance at the minimum compression length exactly balances the gravitational force driving contraction. This is the physical content of what loop quantum cosmology calls the quantum geometry bounce: the substrate's own mechanical properties prevent the singularity.

The bounce is a rebound. The substrate was contracting, collapsing inward under gravity, reached minimum volume, and reversed direction. This reversal is not smooth. A maximum-compression event in a medium with finite modulus produces a shock: a sudden, discontinuous reversal of the compression gradient that propagates outward from the bounce as a compression-rarefaction wave. This is not a metaphor borrowed from fluid mechanics. The substrate obeys the same continuum mechanics as any compressible elastic medium, and the dynamics of the bounce produce exactly the same class of event as compression reversal in any such medium: a shock wave.

This shock wave is the first EM wave. Not a photon yet, the quantized mode structure has not yet condensed, but the condition from which photons crystallize. The bounce shock propagates outward through the rebounding substrate as a spherical compression-rarefaction front. Every point inside the future light cone of the bounce is reached by this wave. There is nothing before this wave in the causal structure of the post-bounce universe.

Immediately at the bounce, the substrate is in a state of maximal geometric stress. Every point in the substrate is simultaneously at maximum curvature. This state is not uniform, and the non-uniformity is critical. Quantum geometric fluctuations, pre-existing topological features carried from the prior contraction cycle, seed non-uniform compression modes across the substrate. Different regions reach slightly different maximum compression densities, at slightly different times, depending on the local substrate geometry inherited from the prior cycle.

These non-uniform modes are the seeds of all structure. The universe does not begin smooth and become lumpy through inflation. It begins with the substrate already carrying geometric inhomogeneities from its prior history, amplified and imprinted by the bounce shock. The amplification occurs because the bounce is a shock and shocks amplify pre-existing perturbations by the same mechanism that acoustic shocks do. Compression at the shock front concentrates any pre-existing density variation. The bounce is, among other things, a perfect amplifier of pre-existing substrate geometry.

The Electromagnetic Field Emerges First

The first thing out of the bounce is not matter. It is field structure. As the substrate rebounds from its maximum-compression state, it does not simply expand uniformly. It oscillates: the rebound drives a compression wave outward, and behind the wave the substrate is in a state of reduced pressure, rarefaction. The wave is not longitudinal only. Because the substrate is three-dimensional and topologically rich, the compression wave immediately generates transverse components: shear waves.

Shear waves in a compressible medium with characteristic tension are precisely what the electromagnetic field is. This identification is exact, not approximate. The stress tensor of a compressible elastic medium under transverse (shear) deformation has the same mathematical structure as the electromagnetic field tensor. The two independent transverse polarization directions of the shear wave correspond exactly to the two polarization states of the electromagnetic field. The dispersion relation of substrate shear waves at low frequencies is:

$$\omega = c_s \cdot k$$

where ω is the angular frequency, k is the wavenumber, and c_s is the shear wave speed of the substrate. This is exactly the dispersion relation of electromagnetic radiation in vacuum, $\omega = ck$. The identification is direct: c, the speed of light, is the shear wave speed of the substrate. This is not a coincidence to be explained away. It is the definition of c. The universality of the speed of light, its independence of source and observer velocity, its role as the maximum signal speed in the universe is a consequence of the substrate's mechanical properties. All signals propagate through the substrate; no signal can propagate faster than the substrate's shear wave speed.

Before any particle condenses, the entire early universe is a sea of EM-mode substrate oscillations. There are no charges yet. There is no distinction between electric field and magnetic field as separate entities, both are projections of the same substrate shear stress tensor onto different spatial directions relative to the wave propagation direction. The electric/magnetic split is observer-dependent and reference-frame-dependent for exactly the reason the substrate model predicts. They are the same deformation mode viewed from different angles.

This EM-field-dominated phase lasts a geometric instant in the post-bounce timeline, not because it is brief in any arbitrary sense, but because the conditions for topological condensation arrive almost immediately. The substrate is cooling, expanding, with energy density falling and as soon as the energy density drops below the threshold for stable knot formation, the first topological structures lock in. That threshold is reached within the first moments of the post-bounce expansion.

How Photons Condense from Substrate Shear Modes

A photon is not a fundamental entity separate from the substrate. A photon is a quantized, self-reinforcing substrate shear packet, a traveling wave mode that has locked into a stable, indivisible unit of substrate oscillation. The discreteness of photon energy, expressed as $E = hf$, is a direct consequence of the substrate's minimum compression length. Oscillation modes shorter than the Planck length cannot exist because the substrate cannot be deformed on scales below its minimum compression length. This discretizes the energy spectrum of substrate shear modes: only integer multiples of a minimum energy quantum are supported. That minimum energy quantum, at a given frequency, is Planck's constant times the frequency. Planck's constant is a substrate mechanical parameter, not an independent fundamental constant.

A photon in this framework is a topologically trivial traveling deformation. It carries no net twist, no net topological charge, and no rest mass. It is a ripple. It cannot stop because the substrate is everywhere and the ripple has no stable resting configuration: it propagates at c (the substrate shear speed) until it transfers its energy to a topological structure, a charged particle, capable of absorbing and storing the deformation. What we call photon absorption is the transfer of a traveling shear packet into a standing topological knot mode. What we call photon emission is the reverse: a topological knot releasing stored strain energy as a traveling shear packet.

The two polarization states of the photon are the two independent transverse shear directions available in three-dimensional space, perpendicular to the propagation direction. The substrate geometry allows exactly two such directions. No more. This is not imposed by hand, it is a consequence of the spatial dimension of the substrate and the transverse character of shear waves. Longitudinal substrate oscillations exist but are distinct: they correspond not to photons but to a different class of mode that I will address in the context of the strong interaction.

The spin-1 character of the photon has an equally direct geometric explanation. A spin-1 particle returns to its original state after a 360° rotation. A substrate shear wave's polarization vector undergoes exactly one full rotation around the propagation axis per wavelength. This is the helicity of a transverse wave: the field vector rotates once per cycle. The photon's spin-1 character is the statement that its internal shear stress vector completes one full revolution per wavelength. This is intrinsic to the geometry of transverse substrate deformation, not imposed as an independent quantum number.

Photons are the earliest stable substrate mode. They do not decay because they carry no topological charge: there is no conserved topological quantity that must be maintained, no knot that must be preserved. A photon is the simplest possible substrate excitation — a pure traveling wave with no internal topological structure. Its stability is the stability of wave propagation itself: a shear wave in a uniform elastic medium propagates indefinitely unless it encounters a topological feature (a charged particle) that can absorb it.

How Electrons Condense

As the post-bounce substrate cools and expands below the energy scale at which stable topological knots can form, the first non-trivial substrate structures lock in: electrons and positrons. These are the earliest massive particles, and they arise not from a symmetry-breaking phase transition in the conventional sense, but from a geometric threshold: below a certain energy density, the substrate can support stable knotted deformations, regions where the topology is non-trivial and cannot be smoothed away.

An electron is a stable, minimum-complexity topological knot in the substrate. More precisely, it is a substrate region that carries a net geometric twist, a topological charge, that cannot be removed by any smooth, continuous deformation of the surrounding substrate. This is the physical meaning of charge conservation: the electron's charge is a topological invariant. You cannot remove it without tearing the substrate. The substrate does not spontaneously develop or annihilate topological boundaries, for the same reason that a knotted string cannot be unknotted without cutting.

The elementary electric charge e is the topological invariant of the simplest stable substrate knot. It is not a dimensionless number chosen arbitrarily, it is the minimum unit of topological charge that the substrate geometry supports, determined by the substrate's mechanical parameters. The fine structure constant $\alpha = e^2/(4\pi\varepsilon_0\hbar c) \approx 1/137$ is the ratio of the electron's topological charge energy to the characteristic energy of substrate shear propagation, a geometric ratio derivable from the substrate parameters, though I defer that calculation to subsequent work.

The electron mass is the rest energy of the knot: the elastic strain energy stored in the substrate deformation around the topological defect. This is not circular. The electron has mass not because it is made of something heavy, but because maintaining the knot geometry against the substrate's tendency to relax requires a constant, localized expenditure of substrate energy. The mass is real, physical, and stored in the geometry of the surrounding substrate, not in any internal structure of the electron itself.

The spin-1/2 character of the electron is the geometric property of the knot that most people find hardest to visualize, and it is the one I am most insistent about explaining geometrically rather than accepting as a quantum-mechanical axiom. A spin-1/2 object requires a 720° rotation, two full turns, to return to its original configuration. This is not an arbitrary quantum rule. It is the property of the simplest non-trivial topological knot in three-dimensional space: the trefoil knot, or more precisely the fundamental braid group element in three spatial dimensions. When such a knot is rotated by 360°, the surrounding substrate deformation does not return to its original configuration, it acquires a phase. A second 360° rotation is required to return the full substrate configuration to its starting state. This double-cover behavior is built into the topology of the minimum-complexity substrate knot. It is not imposed by hand as a quantum number. It emerges.

The positron is the same knot with opposite handedness. Same elastic strain energy (same mass), opposite topological charge (opposite electric charge). When an electron and positron meet, the opposite-handed knots encounter each other, and the substrate finds the lowest-energy resolution: the two opposite twists cancel, unwinding the topology entirely and releasing all stored elastic strain energy as photons (traveling substrate shear waves). This is pair annihilation, a geometrically exact process, not a mysterious quantum event. Two knots of opposite handedness annihilate because their combined topology is trivial: the substrate was unknotted before them and returns to being unknotted after them, with all stored strain energy released as radiation.

The electron is the lightest charged particle because it is the simplest stable knot. Any substrate mode simpler than the electron-knot either carries no topological charge (the photon: topologically trivial) or is unstable (it finds a path to a lower-energy configuration and decays). The electron is the ground state of charged substrate topology. Below it: only the photon. Above it: a spectrum of higher-complexity knots corresponding to the heavier particles.

How Protons Condense

The proton is the most structurally revealing particle in the substrate framework. It is here that the role of prime numbers in determining stable matter becomes explicit and, I will argue, inevitable.

A proton is not a simple knot. It is a composite topological structure: three substrate knot-strands, what the standard model calls quarks, held in a stable geometric configuration by the substrate tension of their mutual interaction. The three strands are not independent. They are topologically linked: a single composite braid in the substrate whose individual strands cannot be separated without unraveling the whole structure. This mutual linkage is confinement, and I will return to it shortly.

The critical question is: why three strands? Why not two, or four, or any other number? The answer is primeness. A composite substrate knot made of n strands is stable if and only if n is prime. A prime-stranded composite knot is topologically irreducible: it cannot be factored into smaller independent sub-knots without breaking the topological coherence of the whole structure. A non-prime-stranded composite, by contrast, can always be decomposed into a product of lower-strand-count composites. The substrate geometry enforces this: a two-strand composite is equivalent to two independent single-strand knots weakly linked. A four-strand composite can be decomposed into two two-strand composites, or a three-strand and a one-strand. These decompositions are geometrically available: the substrate finds them and the structure fragments. A three-strand prime composite has no such decomposition. It is irreducible by topology alone.

This is why mesons, two-quark composites, are not stable: two is not a prime substrate configuration, and the two-strand composite finds a path to decomposition. This is why no stable four-quark particle exists with a macroscopic lifetime: four is not prime. The substrate geometry enforces prime subdivision as the necessary condition for long-term stability. The only stable composite particle count is three, the first prime after the trivial cases of zero and one. Prime subdivision is not numerology. It is the topological condition for irreducibility in a compressible, knotted substrate. The universe is built from prime knots because those are the only knots the substrate cannot undo.

The proton carries a total topological charge of +1 unit of the elementary charge e. This charge is distributed across three strands. The distribution is not arbitrary, it is forced by the constraint that each strand must carry a charge consistent with the overall prime-knot geometry, and the total must be an integer multiple of e (the minimum topological charge unit). Distributing total charge +1 across three strands yields +1/3, +1/3, +1/3 for the up quarks and something more subtle for the down quark configuration (+2/3, +2/3, −1/3). These are not free parameters. They are the unique solutions consistent with the prime-subdivision geometry and integer total charge.

The proton mass is approximately 1836 times the electron mass. In the substrate model, this ratio reflects the elastic strain energy of a three-strand prime-knot composite relative to a single-strand knot. The three-strand composite is more energetically costly not simply because it has three strands, but because those strands are topologically linked: the mutual substrate tension between linked strands contributes additional strain energy beyond the sum of three independent single-strand knots. The factor of 1836 is a geometric property of the three-strand prime composite in a substrate with a specific shear modulus.

The strong force is the substrate tension between the strands of the proton knot. It is not a separate force. It is the substrate resisting the separation of topologically linked strands. Pulling two strands of the proton composite apart requires stretching the substrate between them. As the separation increases, the strain energy in the substrate between the strands increases without bound, there is no finite-energy path to complete separation because complete separation would require an infinite amount of substrate stretching. This is confinement: geometrically exact, requiring no additional mechanism.

The proton is stable indefinitely because its three-strand prime-knot topology has no lower-energy state to decay into. The electron, a one-strand knot, cannot absorb the proton, because the topological charges are not complementary in the right combination. The proton simply persists: an irreducible prime-knot in the substrate, the most stable composite matter structure the geometry allows.

How Neutrons Form

The neutron is the proton's closest structural relative, and its properties are among the most illuminating in the substrate framework because it is the primary example of a non-prime composite, and the substrate's treatment of non-prime composites tells us everything about the nature of nuclear stability.

A neutron is a proton knot (three strands) with an additional bound electron-mode thread woven through its core structure. In standard model language: a neutron contains an up quark, an up quark, and a down quark or equivalently, it is a proton that has absorbed an electron while emitting an antineutrino. In the substrate model, this is a four-strand composite: the three-strand prime knot of the proton with a single additional strand from the electron-mode topology interwoven through it. The total topological charge is zero (the proton's +1 and the electron-mode's −1 cancel), giving the neutron its electrical neutrality.

The neutron in isolation is unstable. The reason is immediate from the primeness condition: four is not prime. A four-strand composite is reducible, it can be decomposed as 4 = 3 + 1. The substrate geometry makes this decomposition available: the four-strand neutron has a finite-energy path to becoming a three-strand proton plus a free one-strand electron-mode, with the topological bookkeeping of the transition carried by what we call the antineutrino. This is beta decay. It is not a probabilistic quantum event without geometric explanation. It is the substrate finding the lower-energy, lower-complexity prime factorization of a non-prime composite. The mean lifetime of the free neutron (≈887 seconds) is the geometric timescale for the substrate to execute this factorization at the energy scale of the neutron mass.

The neutron is stable inside a nucleus for a geometrically distinct reason. Inside the nucleus, the proton and neutron knots are in mutual topological contact: their substrate boundaries are shared and intertwined. The four-strand neutron structure, rather than sitting in isolation, is geometrically entangled with the surrounding three-strand proton structures. The decomposition path, the $4 \rightarrow 3 + 1$ factorization, is blocked because executing it would require separating topological boundaries that are shared with neighboring knots. The neutron cannot decompose without disrupting the surrounding proton structures. In a nucleus with sufficient protons, the energetic cost of nuclear disruption exceeds the energetic gain from neutron decomposition. The neutron remains stable.

This is nuclear binding energy in the substrate framework. It is the energy cost of disrupting the shared topological boundaries between proton and neutron knots in the nucleus. The nuclear force, in the standard model, mediated by pions, is, in the substrate model, the shared-boundary tension between interwoven proton and neutron knot strands. Short-range, strong, and saturable: exactly as the substrate geometry predicts for a tension force between strands that share a common topological boundary of finite size.

The neutron-to-proton mass ratio of approximately 1.001378, the neutron is slightly heavier than the proton, reflects the additional elastic strain energy of the woven fourth strand. The mass difference (about 1.293 MeV) is the strain energy cost of the four-strand configuration relative to the three-strand prime. This is not an accident. The neutron must be heavier than the proton for the substrate model to be self-consistent: if it were lighter, free protons would decay into neutrons, destabilizing all ordinary matter. The substrate geometry ensures the four-strand composite is always slightly heavier than the three-strand prime, by the strain energy of the additional woven strand.

Prime Subdivisions

I want to address the prime number structure of this framework directly and without evasion, because I am aware that invoking prime numbers in a physics context invites the charge of numerology. The charge does not apply here, and I will explain why.

Numerology is the association of numbers with physical phenomena without a causal mechanism. The appearance of prime numbers in this framework is not an association, it is a theorem. In the topology of braids in three-dimensional space, a composite braid is irreducible (cannot be decomposed into smaller independent braids without cutting) if and only if the braid word is not expressible as a product of independent sub-braids. For braids constructed from n strands, this irreducibility condition corresponds directly to the primeness of n. This is a result in braid group theory, not a physical assumption. I am not choosing prime numbers because they seem elegant. The substrate topology selects them because they are the mathematically precise condition for irreducibility in the relevant algebraic structure.

The stable particles we observe correspond to the lowest-complexity irreducible substrate modes:

- **Photon:** Zero net topological charge. Topologically trivial traveling deformation. Stable because there is nothing to unwind. (Strand count: 0 net topological.)
- **Electron:** One-strand knot. The unit topological charge. Stable because one is the unit of the braid group — there is no smaller non-trivial structure. (Strand count: 1.)
- **Proton:** Three-strand prime composite. The first non-trivial prime composite. Stable because three is prime. (Strand count: 3.)

The next prime is five. Five-strand substrate knot composites would correspond to stable particles at energies determined by the five-strand elastic strain energy, substantially higher than the proton mass. These modes are accessible at temperatures present in the early post-bounce universe, and in modern particle accelerators. The heavier quarks and leptons of the standard model: the charm, strange, top, and bottom quarks; the muon and tau leptons are precisely this: higher-prime or higher-complexity substrate knot modes. They are stable at high energy and temperature (when the substrate has sufficient energy density to maintain the higher-complexity knot geometry), and unstable at low energy (when the substrate relaxes to the lowest available prime-knot configuration).

The pattern is clear and general: as the universe cools after the bounce, the substrate undergoes a progressive relaxation from high-complexity modes to low-complexity ones. Each time a higher-complexity knot decays to a lower-complexity prime knot, it releases photons. The early universe is a substrate relaxation cascade, a sequence of topological simplification events, each releasing EM shear radiation. The standard cosmological epochs, the quark-hadron transition, the lepton epoch, the photon epoch, are the observational record of this cascade, seen from the outside.

Within this framework, the four fundamental forces are not independent mechanisms requiring separate gauge theories. They are four aspects of a single underlying geometric reality:

- **Electromagnetic interaction:** Substrate shear stress between topological charge boundaries. The force between electron knots is the stress the substrate develops between two regions of opposite or like twist.
- **Strong interaction:** Substrate tension between strands of a prime-knot composite. Confinement is infinite strand separation energy. Asymptotic freedom is the reduced strand tension at short separations where the substrate deformation overlaps.
- **Weak interaction:** Substrate topology transition. The conversion of a non-prime composite to its prime factorization, with bookkeeping modes (neutrinos) carrying the topological remainder. Every weak decay is a substrate finding a prime factorization.

- **Gravitational interaction:** Large-scale container curvature of the substrate. Discussed in the next section.

Container Geometry

If the substrate is the medium in which all fields and particles exist, what is its large-scale geometry? This question cannot be deferred. The substrate must have a macroscopic shape, a container geometry, and that geometry must be determined by substrate dynamics, not assumed as a pre-given background.

The container is the macroscopic shape of the substrate itself. It is not a separate structure enclosing the substrate from outside. It is the substrate in its large-scale curvature mode, the lowest-frequency oscillation mode of the entire substrate, in which the substrate as a whole deforms. Space does not contain the substrate. Space is the substrate's container mode.

The history of the container geometry follows the bounce directly. Before the bounce, the substrate was in a contracting phase. Its large-scale curvature was positive, a closed geometry, a finite substrate shrinking toward its minimum volume. At the bounce, the substrate reached minimum volume and maximum curvature. After the bounce, the substrate expanded. As it expanded, its curvature decreased. The container geometry evolved from closed (positive curvature, finite volume) through the bounce, outward toward flatness (zero curvature, infinite effective volume) as the substrate relaxed.

The observed flatness of the universe, the near-perfect spatial flatness measured by CMB observations, is not a fine-tuning problem requiring inflation to solve. It is a consequence of the substrate's bounce-driven relaxation dynamics. A substrate compressed to maximum curvature and then released does not immediately settle into its equilibrium configuration. Like a compressed spring released past its equilibrium, it overshoots. The substrate expands through flatness, overshoots slightly into negative curvature, and oscillates (on cosmological timescales) toward equilibrium. The CMB-era flatness corresponds to the substrate near its first zero-crossing after the bounce: close to flat because the bounce compressed it to maximum curvature and the rebound has carried it very nearly to the flat equilibrium configuration. No inflation required.

General relativity's field equations are the continuum mechanics equations of the substrate, linearized at low curvature. Einstein's field equations describe how the substrate container geometry responds to the distribution of topological knot energy (mass-energy) within it. The stress-energy tensor is the substrate's knot-energy distribution. The Einstein tensor is the substrate's curvature response. The equations are correct as a description of substrate container dynamics; they are incomplete as an explanation because they treat the substrate as an abstract geometric manifold rather than a physical medium with specific mechanical properties.

Dark energy is the residual elastic tension of the substrate container. The substrate was compressed to maximum curvature at the bounce and has since expanded far beyond its equilibrium configuration. A medium expanded past equilibrium exerts a restoring force, but the direction of this force depends on the equilibrium state. If the substrate has expanded past its equilibrium volume, the elastic restoring force acts outward, accelerating further expansion. This is what we observe as the cosmological acceleration. The cosmological constant Λ is not a vacuum energy density added by hand. It is the substrate's elastic modulus at cosmological scale, describing the restoring force of a substrate expanded beyond its equilibrium configuration. Its value is set by the substrate's mechanical parameters, not by the vacuum energy density of quantum field theory. This resolves the cosmological constant problem: the discrepancy between the observed value of Λ and the quantum field theory prediction is not a puzzle, because they are measuring different things. QFT vacuum energy is a description of substrate zero-point shear modes. The cosmological constant is the substrate's large-scale elastic modulus. They have no reason to be equal.

The CMB — Fossil Record of the First Substrate Wave

The Cosmic Microwave Background is the most important observational data set for this framework. Not because it confirms inflation, it does not, at least not cleanly, and the B-mode tensor mode signature that would confirm inflation has not been unambiguously detected, but because it is exactly, feature by feature, what a first substrate oscillation wave would produce. I will go through the CMB's observed properties one by one and show that each is a direct geometric consequence of the bounce wave picture.

Here is what happened. At the bounce, the substrate released a compression-rarefaction shock wave: a spherical shell of EM-mode substrate oscillation propagating outward at c. This is the photon epoch. As the post-bounce substrate cooled, the relaxation cascade proceeded: quarks condensed, protons and electrons formed. The EM-mode oscillations, the first substrate wave, continued propagating but were now scattered by the free charged substrate knots (electrons). The early universe was opaque to EM radiation not because it was a hot gas in equilibrium, but because free electron knots scatter EM shear waves efficiently: a traveling shear packet (photon) interacts with any topological charge boundary (electron) it encounters.

At recombination, approximately 380,000 years after the bounce, at a temperature of about 3000 K, electrons were captured into hydrogen atoms. Their topological charge became bound into electrically neutral composites. Neutral composites interact with EM shear waves far more weakly than free charge boundaries do. The universe became transparent. The original bounce wave, which had been propagating and scattering since the first moments, was suddenly free to propagate without further scattering. This is the surface of last scattering, and what we observe as the CMB is the original bounce wave at this surface, redshifted by the subsequent expansion of the substrate container.

The perfect blackbody spectrum of the CMB, the most precisely measured blackbody spectrum in physics, with deviations less than one part in 10^4, is not evidence of thermal equilibrium in the conventional sense. A gas in thermal equilibrium produces a Planck spectrum because every mode is populated according to the Boltzmann distribution. The bounce wave produces a Planck spectrum for a different and more fundamental reason: a substrate oscillation wave produced at a single event (the bounce) excites every mode simultaneously, with mode energies determined by the substrate's mechanical dispersion relation and the geometry of the bounce event. The result is a geometric Planck distribution: every mode is excited at the same moment, with amplitudes determined by the substrate's shear modulus and the symmetry of the bounce. This is a Planck distribution not because of thermalization but because of geometric mode excitation. The perfection of the CMB blackbody spectrum is evidence not for a hot plasma in equilibrium, but for a single-event geometric excitation of a substrate: the bounce.

The CMB temperature anisotropies, temperature fluctuations at the level of one part in 10^5 are the imprint of pre-existing geometric non-uniformities in the substrate at the moment of the bounce. These are not generated by inflation. They are the topological memory of the substrate from its prior contraction cycle, amplified by the bounce shock. The substrate carries its prior history through its geometry. The anisotropies are that history, stamped into the first substrate wave at the bounce, and preserved through recombination, now readable as temperature fluctuations in the CMB sky.

The acoustic peaks in the CMB power spectrum, the series of peaks in the angular power spectrum at specific multipole moments are the standing wave modes of the substrate container during the recombination era. The bounce wave traveled outward from the bounce, and the substrate container (with its finite curvature at that epoch) acted as a resonant cavity. The bounce wave set up standing wave modes in this cavity: the resonant harmonics of the substrate container at the acoustic horizon scale. When recombination occurred, these standing wave modes were frozen in, the pattern of compressions and rarefactions was imprinted in the photon-baryon fluid at the moment it decoupled. The peak positions in the CMB power spectrum are the integer harmonics of the acoustic horizon: the distance the bounce wave traveled from the bounce to recombination, divided by integers 1, 2, 3, and so on. This is why the peaks occur at specific, predictable angular scales. They are the geometric harmonics of the substrate container, ringing with the original bounce wave. Their positions are determined entirely by the substrate's sound speed during the recombination era and the geometry of the container at that time.

The CMB is not an afterglow of the Big Bang in the sense of a hot gas cooling to equilibrium. It is the direct, observable, redshifted record of the first substrate oscillation wave. It is the bounce, still propagating, 13.8 billion years later, now stretched to microwave wavelengths by the expansion of the substrate container. Every feature of the CMB, its blackbody perfection, its temperature of 2.725 K, its anisotropy spectrum, its acoustic peaks, its polarization pattern, is a geometric record of substrate dynamics during and immediately after the bounce. The CMB is the most detailed observational window into the bounce that exists. We have been looking at the bounce itself all along.

The Complete Sequence

The complete causal chain from the bounce to the CMB is as follows. Each step is a direct geometric consequence of the preceding step, with no gaps and no additional physical assumptions beyond the three substrate parameters.

Bounce. The substrate reaches minimum volume and maximum compression. Its compression modulus prevents further contraction. A geometric shock is initiated: a compression-rarefaction wave propagates outward from the point of maximum compression into the rebounding substrate. The pre-bounce topological memory of the substrate is preserved and amplified by the shock.

EM Field Emerges. As the substrate rebounds, shear waves propagate outward alongside the compression wave. Substrate shear waves are the electromagnetic field. The shear wave speed c is the speed of light. No charges exist yet; no distinction between E and B fields has meaning. The universe is a pure EM-mode substrate oscillation.

Photons Condense. The substrate's minimum compression length discretizes the shear wave energy spectrum. Quantized, self-reinforcing shear packets, photons, lock in as the earliest stable substrate mode. Topologically trivial (no charge, no mass), they propagate at c indefinitely. Spin-1 character is the geometry of transverse shear. Two polarization states are the two transverse shear directions.

Electron and Positron Condensation. As the substrate cools, the energy density drops below the threshold for stable knot formation. The first topologically non-trivial substrate modes lock in: one-strand knots with net topological charge. The electron (left-handed knot) and positron (right-handed knot) form. Charge is a topological invariant. Mass is elastic strain energy of the knot. Spin-1/2 is the geometric property of the minimum-complexity non-trivial knot.

Quark Condensation. Higher-prime substrate knot modes condense as the substrate continues cooling. Three-strand prime composite modes form, with strand charges forced to fractional values (±1/3, ±2/3) by the prime-subdivision geometry. These are quarks: the strands of the proton and neutron knots.

Proton Condensation. Three-strand prime composites stabilize. Substrate tension between the linked strands (the strong force) binds them in an irreducible prime knot. Proton mass is the elastic strain energy of the three-strand composite, approximately 1836 times the electron knot energy. Proton stability is guaranteed by the primeness of the three-strand topology: no lower-energy factorization is available.

Neutron Formation. Four-strand composites form: three-strand proton knots with an additional electron-mode strand woven through. Electrically neutral (charges cancel). Unstable in isolation because 4 is not prime, the substrate finds the 3+1 factorization (beta decay) on a timescale of ≈887 seconds. The antineutrino carries topological bookkeeping of the transition.

Nuclear Formation. In regions of sufficient density, proton and neutron knots come into mutual topological contact. Shared-boundary entanglement stabilizes the neutron (blocks the 3+1 factorization) and binds protons and neutrons into nuclei. The nuclear force is shared-strand-boundary tension. Nuclear binding energy is the geometric cost of disrupting shared boundaries. Stable nuclei form where topological entanglement provides sufficient stabilization energy.

Container Geometry. The large-scale substrate curvature (the container) relaxes from its maximum-curvature bounce state. GR field equations are the substrate continuum mechanics at low curvature: correct as description, incomplete as explanation. The universe relaxes toward spatial flatness by bounce-driven dynamics, without inflation. Dark energy is the substrate's elastic tension in its current super-equilibrium expanded state; the cosmological constant is the substrate's large-scale elastic modulus.

Recombination. Approximately 380,000 years after the bounce, the substrate temperature drops below the ionization energy of hydrogen. Free electron knots are captured into neutral hydrogen atoms. The free topological charge boundaries that had been scattering the EM shear wave (the original bounce wave) disappear. The universe becomes transparent. The bounce wave decouples from matter.

CMB. The original bounce wave propagates freely after decoupling. Its current energy density, redshifted by the expansion of the substrate container, corresponds to a temperature of 2.725 K. Its blackbody perfection reflects geometric mode excitation at the bounce. Its anisotropies are the substrate's topological memory of its pre-bounce geometry, stamped into the first wave at the bounce. Its acoustic peaks are the geometric harmonics of the substrate container. The CMB is the direct, observable record of the bounce itself.

Conclusion

The bounce is not a metaphor. It is not a mathematical convenience. It is the physical consequence of the substrate’s global geometry. The universe is a continuous spherical substrate with finite elastic tension and a maximum expansion strain. As the post-bounce universe expands, the substrate container stretches. This stretch is not unlimited. A continuous sphere cannot expand beyond the strain its geometry permits without tearing, and tearing is forbidden: the substrate has no seams, no boundaries, and no allowable topological transition that would permit rupture. The geometry enforces a maximum radius.

When the universe reaches this maximum expansion, the substrate is at its limit. Any further outward motion would require a discontinuity in the substrate, a tear, and the substrate does not tear. At this point the container tension reverses the sign of acceleration. Outward velocity may still be positive, but the acceleration is everywhere inward. The substrate's elastic geometry forces all matter, radiation, and substrate modes back toward the center. This reversal is the bounce.

The bounce is the turning point of the universe's expansion: the moment at which the spherical substrate reaches its maximum allowable strain and the only physically permitted evolution is inward motion. The universe does not bounce because it was collapsing. It bounces because it expanded as far as a continuous spherical substrate can expand. The geometry of the substrate, its continuity, its finite tension, and its prohibition against tearing, determines the bounce uniquely. The next contraction, the next cycle, and the next expansion all follow from this single geometric fact: a sphere with no seams cannot expand forever.

The freeze–bounce of the cosmic substrate is the same physical geometry seen in a neutron star hitting its compression limit before a supernova: a continuous medium reaches maximum strain, can no longer compress, and the only permitted evolution is an inward recoil followed by explosive release.