

Reaction Pathways as Geometry Routes: Tension Windows, Alignment Constraints, and the Mechanical Logic of Chemical Selectivity

Dustin Lee

Abstract

Chemical reaction pathways are not abstract energy diagrams. They are geometry routes. A molecule takes a specific pathway because its electron domains, bond angles, and mechanical strain state physically route it there. The familiar reaction coordinate diagram is not the cause of anything — it is a downstream record of geometry decisions the molecule has already made. This paper builds that argument in ten sequential sections. Section 1 establishes that pathways are geometry sequences, not energy slopes. Section 2 defines the **tension window** — the range of mechanical strain within which a geometry transition is accessible — and explains why it gates pathway entry. Section 3 treats geometry alignment as the actual gating event: the simultaneous satisfaction of multiple spatial constraints required for a bond-forming or bond-breaking event to proceed. Section 4 explains electron-domain routing as the physical basis of arrow-pushing. Section 5 analyzes how competing pathways fight over the same molecular population through overlapping tension windows and competing alignment events. Section 6 examines how catalysts reshape geometry routes rather than simply lowering activation barriers. Section 7 catalogs the mechanical failure modes of reaction pathways. Section 8 explains side reactions as geometry drift events, not random errors. Section 9 synthesizes these principles into a practical framework for reaction control. Section 10 presents a systematic procedure for diagnosing and restoring failed geometry routes. Throughout, mechanism is treated as primary. Energy diagrams are treated as summaries of geometry decisions already made.

1. Pathways Are Geometry Routes, Not Energy Slopes

The standard picture of a chemical reaction shows an energy diagram: reactants on the left, a hill in the middle, products on the right. The molecule, as the diagram implies, rolls up the hill and then down the other side. This picture is not wrong, but it is incomplete in a way that matters enormously. The energy profile is not the *cause* of the pathway. It is the *record* of it. The molecule does not move because of an abstract energy gradient pulling it forward. It moves because specific atomic positions become mechanically favorable and others become mechanically blocked. The energy numbers follow from the geometry. If you want to understand why a reaction takes the pathway it takes, you have to start with the geometry, not with the energy.

A **reaction pathway**, properly understood, is a sequence of geometry states the molecule passes through. Each state in that sequence has specific bond lengths, bond angles, electron-domain orientations, and strain distributions. The pathway is defined by which geometry states are accessible in sequence — which configurations can be reached from the starting configuration, and in what order. Some pathways thread smoothly through a series of geometrically connected states. Others are blocked not because the energy is abstractly high, but because the transition-state geometry requires bond angles or domain orientations that the molecule's current physical configuration cannot reach. The geometry is the actual constraint. The energy surface describes what happens after geometry is fixed, not why geometry goes where it goes.

VSEPR geometry — the Valence Shell Electron Pair Repulsion model — is the baseline logic here. Electron domains repel each other and find minimum-strain configurations. A molecule in a given bonding arrangement will adopt the geometry that minimizes the repulsion between all of its electron domains. The pathway through chemical space is the sequence of VSEPR-valid configurations the molecule can thread through, starting from reactant geometry and ending at product geometry. Each intermediate geometry along the

route must itself be a mechanically accessible arrangement — one that the electron domains can occupy without requiring impossible compressions or expansions of bond angles. The pathway is not a smooth slide down an energy slope. It is a specific physical route through a landscape of geometry states, some of which connect to each other and some of which do not.

The clearest illustration of this is the **SN₂ displacement reaction**. In SN₂, a nucleophile attacks a carbon center bearing a leaving group, and the leaving group departs as the new bond forms. The energy diagram shows a single-step profile with a transition state at the top. But that diagram says nothing about why the nucleophile must approach from the back side of the C–X bond. The reason is purely geometric. The carbon's sigma* antibonding orbital — the orbital the nucleophile must donate into to break the C–X bond — has its largest accessible lobe on the back side of the bond, pointing directly away from the leaving group. The nucleophile's lone pair must align with that lobe. If the nucleophile approaches from the front, there is no orbital overlap of the right symmetry to initiate bond formation. The geometry of the antibonding orbital locks the approach geometry.

Furthermore, as the nucleophile approaches and the new bond begins to form, the carbon must pass through a **trigonal bipyramidal transition state** in which the three remaining substituents are arranged in a flat plane perpendicular to the nucleophile–C–leaving-group axis. Carbon inverts through that planar geometry — the Walden inversion — and emerges on the other side with its configuration flipped. This inversion is not an abstract activation energy event. It is a specific geometry transition: three bond angles that were tetrahedral must flatten to 120° simultaneously as the two axial bonds stretch and compress. The energy barrier is the cost of achieving and passing through that geometry. The geometry route is: nucleophile aligns with sigma* backside lobe, carbon flattens to trigonal bipyramidal, leaving group departs, carbon reforms tetrahedral with inverted configuration. That is the pathway. The energy diagram summarizes it, but the geometry is what defines it.

A pathway that is **geometry-blocked** is not simply a pathway with high activation energy. It is a pathway where the required transition-state geometry is physically unreachable from

the starting configuration. A neopentyl substrate in an $\text{S}_{\text{N}}2$ reaction is not just "sterically hindered" in a vague sense. The three methyl groups on the beta carbon physically prevent the nucleophile's lone pair from achieving the linear alignment geometry with the sigma* orbital that $\text{S}_{\text{N}}2$ requires. The approach geometry is blocked before any bond-breaking chemistry can begin. The molecule cannot thread the nucleophile into the correct geometry, so the reaction does not proceed by $\text{S}_{\text{N}}2$ regardless of how reactive the nucleophile is or how good the leaving group is. Understanding this is the difference between memorizing a rule ("neopentyl doesn't do $\text{S}_{\text{N}}2$ ") and understanding the machine.

Mechanical Summary — Section 1

Reaction pathways are sequences of geometry states. The energy profile is the cost of passing through those states. The geometry determines what states are accessible. Geometry-blocked pathways are not high-energy pathways — they are pathways where the required transition-state geometry cannot be reached from the starting configuration. Start with geometry; let energy follow.

2. Tension Windows: What They Are and Why They Gate Pathways

Every reaction pathway has a geometry requirement at the transition state. The molecule must thread through a specific arrangement of atoms to get from reactant to product. But the molecule does not start at that transition-state geometry — it starts at the reactant geometry, which is different. The question of whether a pathway is accessible is therefore not just about the transition-state geometry in isolation. It is about whether the molecule's current mechanical state allows it to reach that geometry. That is what a **tension window** describes. A tension window is the range of mechanical strain states a molecule can occupy while still allowing a specific geometry transition to occur.

Tension in a molecule has two primary components. The first is **bond-length strain**: compression or extension of bonds away from their equilibrium lengths. Bonds have a natural resting length determined by the sizes and electronegativity of the bonded atoms. When a bond is compressed shorter than that length, it stores energy and resists further compression. When it is stretched longer, it weakens toward dissociation. The second component is **angle strain**: deviation of bond angles from their VSEPR-optimal values. A carbon with four single bonds wants angles near 109.5° . Force those angles smaller — as in cyclopropane, where angles are constrained to 60° — and you store significant mechanical strain. Force them larger, and you create a different kind of destabilization. Both bond-length strain and angle strain create mechanical resistance, and both factor into whether a molecule can thread through a particular transition-state geometry.

The tension window is the envelope of strain values — in both dimensions — within which the geometry transition is mechanically accessible. This window has both a lower bound and an upper bound. The lower bound exists because a completely relaxed molecule lacks the pre-organized geometry needed to reach the transition state. Reaching a transition state requires some degree of geometric distortion from the ground state. If the molecule is fully relaxed with no strain, all of the distortion must be introduced from scratch by thermal energy alone. Some geometry transitions, particularly those requiring simultaneous distortion of multiple bonds and angles, cannot be reliably achieved this way — the probability of all the required distortions occurring simultaneously, without any pre-organization from molecular strain, is too low. A small amount of pre-existing strain, oriented in the direction of the required transition-state distortion, provides the geometric head start that makes the transition accessible.

The upper bound of the tension window exists because excessive strain distorts the electron-domain arrangement so severely that the transition-state geometry becomes inaccessible from the opposite direction. A molecule that is too strained has its electron domains mis-oriented. Its bond angles are already compressed or expanded past the point where the additional distortions required by the transition state can be accommodated.

The molecule is deformed away from the geometry it needs, not toward it. The window closes because the molecule is already past the geometry it needs to be approaching from.

A ring system provides the clearest concrete example. A six-membered ring in the **chair conformation** has a specific tension state: the ring bonds are close to their equilibrium lengths, the bond angles are near 109.5° , and the ring is in a well-defined geometry that places axial and equatorial positions at fixed, predictable orientations. A nucleophile attacking a carbonyl on the ring from the axial direction must approach along a trajectory that satisfies the Bürgi–Dunitz angle — roughly 107° relative to the C=O bond — and that trajectory is well-aligned with the π^* orbital of the carbonyl when the ring is in chair. The ring in chair conformation is inside the tension window for axial nucleophilic addition. Now introduce substituents that flatten the ring toward a **boat conformation**. The ring's geometry changes: axial and equatorial positions are no longer clearly defined, the approach trajectory for axial attack is now poorly aligned with the π^* orbital, and the geometry of the ring can no longer guide the nucleophile into the correct approach angle. The tension window for axial addition has collapsed. The product distribution changes accordingly — attack occurs from a different trajectory, or a different reaction pathway dominates.

Solvent, temperature, and substituents all shift a molecule into or out of a tension window. Solvent affects the tension state of the molecule by interacting with its electron domains — a polar protic solvent that hydrogen-bonds to an ester oxygen is physically compressing the geometry of that oxygen's electron-domain arrangement, which changes the local bond angles and partially activates the carbonyl for attack. Temperature shifts the population of molecules occupying strained conformations: higher temperature populates geometries that are higher in energy, potentially pushing more molecules into the tension window or pushing them out of it. Substituents introduce or relieve angle strain and torsional strain directly, reshaping the tension state of the molecule at every conformationally connected site.

The tension window is a **gating concept**. The pathway is only open when the molecule is inside the window. This is why identical reactions run at different temperatures or in different solvents give different product distributions. The tension state of the molecular population changes with conditions, and the tension window for each competing pathway opens or closes accordingly. A reaction that gives substitution at low temperature and elimination at high temperature is not changing which pathway is energetically preferred in an abstract sense. It is changing the distribution of molecular tension states, opening the elimination tension window and closing the substitution window. The molecules are doing what their geometry tells them to do. The conditions determine which geometry the molecules are in.

3. Geometry Alignment: The Actual Gate

Geometry alignment is the specific, simultaneous satisfaction of multiple spatial constraints required for a reaction to proceed. Being inside the tension window is necessary but not sufficient. The molecule must also achieve, at the same instant, a particular arrangement of its reacting atoms and orbitals in three-dimensional space. Alignment is harder to achieve than tension-window entry because it requires not just a range of acceptable strain states but a precise geometric configuration — multiple constraints satisfied at once, not sequentially.

Geometry alignment for a bond-forming or bond-breaking event involves at minimum four simultaneous constraints. First, the **donor orbital** must point directly toward the acceptor orbital. Not approximately — the orbital overlap that drives bond formation drops off sharply with angular deviation from the ideal approach trajectory. Second, the bond angle at the reacting center must be within tolerance. The reacting atom must be in a geometry that can accommodate the incoming bond without requiring an impossible local distortion. Third, the leaving group or displaced fragment must be oriented to exit without steric collision with the rest of the molecule. A leaving group that departs directly into a wall of

substituents creates a geometry conflict that blocks the transition state even if every other constraint is satisfied. Fourth, the overall molecular geometry must be consistent with the transition-state geometry — not just locally at the reacting atoms, but globally across the molecule, because the transition state puts constraints on dihedral angles and ring conformations that can propagate far from the reaction center. All four must be satisfied simultaneously. The reaction does not proceed if even one is missing.

Molecules are not rigid. They are continuously sampling geometry space through rotation around single bonds, ring puckering, bond vibration, and overall conformational dynamics. At any given moment, a molecule occupies one point in this geometry space. Alignment happens only when that point corresponds to a conformation that satisfies all four constraints at once. The reaction rate is therefore partly a function of the fraction of time the molecule spends in aligned conformations, and partly a function of how quickly it can re-enter aligned conformations after leaving them. This is not a correction to collision theory or transition-state theory — it is what those theories are actually computing when they describe conformational populations and pre-exponential factors.

Orbital overlap is a geometry concept. It is not an energy concept, though it has energy consequences. Sigma bond formation requires head-on overlap: the two contributing orbitals must point directly at each other. Pi bond formation requires parallel p-orbital alignment: the two p orbitals must be parallel and their lobes must be adjacent. Lone-pair donation into an empty or antibonding orbital requires the lone pair to point directly at the acceptor orbital lobe. These are geometry requirements. The energy gain from achieving overlap is the reward for satisfying the geometry. You cannot achieve the energy gain without first satisfying the geometry.

The **E2 elimination reaction** is the definitive illustration of alignment as the actual gate. In E2, a base removes a proton from one carbon while the C-leaving-group bond breaks simultaneously, forming a pi bond between the two carbons. For this to work, the H being removed, the carbon bearing it, the carbon bearing the leaving group, and the leaving group itself must all lie in the same plane, with the H-C-C-LG **dihedral angle near 180°**. This is

the anti-periplanar requirement. The reason is orbital geometry: the sigma bond breaking at C-H and the sigma bond breaking at C-LG must both donate electron density into the new pi bond. That requires the C-H sigma bond and the C-LG sigma bond to be parallel to each other and parallel to the forming pi bond's p orbitals. Parallel means anti-periplanar when the two bonds are on adjacent carbons. The 180° dihedral is not a preference or a tendency. It is a hard geometric requirement for the simultaneous orbital overlap that E2 demands.

If the molecule cannot adopt an anti-periplanar conformation with respect to the H and LG in question, E2 does not proceed at that site regardless of base strength. In rigid ring systems where the relevant H and LG are locked in gauche or syn relationships, E2 is shut off entirely. The base is strong enough. The leaving group is reactive enough. But the molecule cannot achieve the geometry, and the reaction does not happen. This is a geometry failure, not an energy failure. The distinction matters because the fix is geometric — redesign the molecule to allow the anti-periplanar conformation — not energetic. Adding more reactive base does not fix a geometry problem.

Anti-periplanar elimination also demonstrates why alignment is a multi-constraint problem. It is not enough to have the H-C-C-LG dihedral near 180° . The base must approach the H from a direction consistent with its role as a proton acceptor while the leaving group departs on the opposite side, and the forming pi bond must have space to develop without colliding with adjacent substituents. If the molecule threads through the anti-periplanar geometry but one of the adjacent substituents blocks the forming pi bond, the elimination pathway still fails. All constraints must be satisfied at once.

4. Electron-Domain Routing: How Charge Moves Through Geometry

Electron domains are not uniformly distributed blobs of charge. They occupy specific regions of space determined by the geometry of the atoms around them. A bonding pair between two carbons occupies the space between those carbons along the internuclear axis.

A lone pair on a nitrogen in a pyramidal amine occupies a specific lobe of space pointing away from the three N-H or N-C bonds. A pi bond occupies two lobes of space above and below the plane of the bonded atoms. These are physical locations. When a reaction proceeds and electron density moves, it moves from one of these physical locations to another. The route that electron density takes through the molecule during a reaction is what **electron-domain routing** describes — and that route is determined entirely by the geometry of the electron domains the density must pass through.

Arrow-pushing in organic chemistry is a notation for electron-domain routing. Each curved arrow represents electron density moving from a filled domain (lone pair, bonding orbital, pi bond) to an adjacent or overlapping domain (empty orbital, antibonding orbital, developing bond). That movement is only possible when the source domain and the accepting domain are spatially adjacent and appropriately oriented. An arrow drawn in a mechanism is not a bookkeeping abstraction — it is a geometry claim. It says that these two electron domains are oriented such that density can move from one to the other. If the geometry does not support that orientation, the arrow cannot be drawn, and the mechanistic step does not occur.

Electron-domain routing is directional. A lone pair on nitrogen can donate into a pi* antibonding orbital only if the nitrogen lone pair is geometrically oriented to overlap with the pi* lobe. This requires the lone pair to be in a geometry that places it parallel to the pi system — pointing into the same orbital space that the pi* lobe occupies. In an aromatic ring where a nitrogen's lone pair is incorporated into the pi system (as in pyrrole), the lone pair is already committed to that routing: it is in a p orbital parallel to the ring pi system, and it routes continuously through the ring. It is not available for donation to an external electrophile in the same way an aliphatic amine's lone pair is. The aliphatic amine's lone pair is in an sp³-hybridized orbital pointing away from the molecule — it is geometrically aimed outward, available for external donation. The geometry of the lone pair's orbital determines the direction of its routing. Different geometries, different routing behavior, completely different chemistry.

In a **conjugated diene**, pi electron density is routed across all four carbons continuously. The reason is purely geometric: the four p orbitals of the four carbons are all parallel, their lobes are adjacent, and they form a continuous overlap pathway. Electron density placed anywhere in this pi system can move across all four carbons because every step of the route is geometrically allowed. The overlap is continuous. Now force the diene out of planarity — twist one end of the diene so that the p orbitals at C₂ and C₃ are no longer parallel. The overlap pathway breaks at that junction. Electron density can still route through C₁–C₂ and through C₃–C₄, but it cannot route across C₂–C₃ because the orbitals no longer overlap. The conjugation breaks. This is not a change in bond orders or resonance structures in an abstract sense — it is a physical interruption of the geometry route through which electron density was flowing. The route closes when the geometry closes it.

Hyperconjugation works the same way. C–H sigma bond electrons can route into an adjacent empty or antibonding orbital when the C–H bond is aligned appropriately — anti-periplanar or synperiplanar — with the acceptor orbital. The C–H sigma bond occupies a specific lobe of space. The adjacent empty orbital (say, a p orbital on a neighboring carbocation) also occupies a specific lobe of space. When those two lobes are parallel and adjacent, their overlap is non-zero, and electron density from the C–H bond can route into the empty orbital. When the C–H bond is gauche to the empty orbital, the lobes are not parallel, the overlap is negligible, and routing does not occur. Hyperconjugation is not a vague stabilizing "interaction" — it is electron-density routing through a geometry-determined pathway, exactly like conjugation, except through a sigma bond instead of a pi bond.

The practical consequence of this is direct and important: **designing where electrons go means designing the geometry** that the routing must pass through. If you want a lone pair to donate into a specific orbital, you arrange the molecule so that the lone pair is geometrically aimed at that orbital. If you want to block a routing pathway, you introduce a geometry constraint that closes the overlap between the relevant domains. Every protecting group, every directing group, every conformational lock is, at the electron-

domain level, a geometry intervention in the routing network. The chemistry follows the geometry route. Change the geometry, and the electrons route differently.

5. Branch-Window Competition: How Pathways Fight Over the Same Molecule

Most molecules under realistic reaction conditions are not committed to a single pathway. They occupy a thermal distribution of geometry states, and multiple pathways may be accessible simultaneously — or nearly so. The question of which product forms is not decided by any single molecular event. It is decided by a continuous competition between pathways for the same molecular population. Understanding that competition requires understanding how the tension windows and alignment requirements of competing pathways interact with the geometry-state distribution of the molecular population. This is the branch-window problem.

A **branch window** is the region of molecular geometry space where two or more pathway tension windows overlap. Inside a branch window, the molecule can satisfy the entry requirements of more than one pathway at once. It is not committed. What determines which pathway it takes is the relative frequency of complete alignment events for each competing route. Pathway A requires alignment set A — a specific combination of dihedral angles, approach trajectory, and orbital orientation. Pathway B requires alignment set B — a different combination. If the molecule samples both conformational families at comparable rates, both pathways receive a share of the molecular population proportional to their alignment-event frequencies. The product ratio is the ratio of those frequencies, not the ratio of abstract activation energies, though those numbers are related. The geometry determines the frequencies, and the frequencies determine the product distribution.

Temperature enters this picture at the conformational sampling level. Higher temperature does not simply speed up all reactions proportionally. It changes which conformational

states are thermally populated. Some conformational states are low in energy and densely populated at all temperatures. Others are significantly strained and only meaningfully populated at high temperature. If pathway B requires a strained conformation for alignment and pathway A does not, then as temperature rises, the population of molecules capable of pathway-B alignment events increases relative to pathway A. The product ratio shifts. This is not because pathway B has a lower activation energy than pathway A in the Arrhenius sense — the Arrhenius equation is computing the same thing in mathematical form, but the physical interpretation is conformational population dynamics, not an abstract thermodynamic preference.

The competition between **SN₂ substitution** and **E₂ elimination** is the canonical branch-window problem in organic chemistry. Both pathways want the same electrophilic carbon and use the same leaving group. SN₂ requires the nucleophile's lone pair to approach from the back side of the C–LG bond, aligned with the sigma* orbital of that bond. E₂ requires the base to approach the beta hydrogen from a direction consistent with deprotonation while the H–C–C–LG system is anti-periplanar. Both requirements can be satisfied by the same molecule in different conformations. The branch window is the conformational space where both pathways are energetically and geometrically accessible.

What closes one branch and opens the other is geometry engineering. **Steric bulk around the reaction center** narrows the SN₂ alignment window: as the carbon becomes more substituted (primary to secondary to tertiary), the approach trajectory for backside attack becomes increasingly obstructed by the substituents on that carbon. The nucleophile cannot physically reach the geometry required for SN₂ alignment. The SN₂ branch window closes. Simultaneously, the E₂ branch window may remain open or widen: the anti-periplanar beta-H–C–C–LG arrangement is not blocked by substitution at the alpha carbon, and a bulky base may actually prefer the accessible proton on the less hindered beta carbon, favoring a Hofmann-type elimination product. The substituents at the alpha carbon do not affect the E₂ alignment geometry in the same way they affect SN₂. The two pathways have different geometry sensitivity to the same structural change.

A good anti-periplanar hydrogen widens the E2 alignment window. If the molecule has a beta hydrogen that is easily accessible in an anti-periplanar arrangement with the leaving group — because the ring geometry or chain conformation places it there preferentially — then the E2 alignment event is highly probable, and the E2 pathway captures a large fraction of the molecular population. A molecule where all beta hydrogens are locked in gauche or syn relationships to the leaving group (as in certain rigid bicyclic systems) cannot achieve E2 alignment at all, and the E2 branch window is geometrically closed regardless of how strong the base is.

Selectivity is branch-window management. The chemist's job in a competitive-pathway situation is to close the windows they do not want and open the ones they do. This is not vague optimization — it is targeted geometry intervention. Identify which conformation allows the competing pathway's alignment event. Introduce a constraint that closes that conformation. Change the solvent to shift the tension window. Adjust temperature to shift conformational populations. Use a bulky reagent that cannot fit the geometry of the competing alignment. Each of these is a geometry decision that changes which branch windows are open. The product distribution reflects the geometry decisions you made, not an inscrutable chemical preference.

6. Catalyst-Driven Geometry Reshaping

A catalyst does not lower activation energy in the abstract. That phrase, repeated endlessly in textbooks, obscures what is actually happening. A catalyst changes the geometry of the transition state, or it changes the geometry of the substrate so that the transition state is easier to reach, or it constrains the substrate geometry so that only one geometry route is accessible. These are three distinct mechanisms, and all of them are geometry interventions, not energy interventions. The energy change is the consequence. The geometry is the cause.

Lewis acid catalysts work by coordinating to a lone pair on the substrate — most commonly on oxygen, nitrogen, or a halogen — and pulling electron density away from that atom. This changes the bond lengths and bond angles around the coordinated atom in a specific way. The bond to the Lewis acid forms and shortens. The bonds from the coordinated atom to the rest of the molecule lengthen and weaken, because electron density that was shared in those bonds has been pulled toward the Lewis acid. In a carbonyl compound, Lewis acid coordination to the carbonyl oxygen weakens the C=O pi bond by pulling pi electron density toward the oxygen and the Lewis acid. The carbon becomes more electron-deficient and, critically, the geometry of the carbonyl changes: the C=O bond length increases, and the carbon's electrophilicity increases because its p orbital is more available for nucleophilic attack. The substrate is now pre-strained in the direction of the transition state for nucleophilic addition. The geometry has been reshaped. The tension window for nucleophilic addition is now accessible at lower temperature or with a less reactive nucleophile — not because the energy barrier is abstractly lower, but because the substrate geometry is already partway through the required distortion.

Enzyme active sites are three-dimensional geometry templates. The active site is a pocket of precise shape and charge distribution, evolved to bind the substrate and physically position it in a conformation that satisfies all the alignment requirements for the target transition state simultaneously. The enzyme does not provide activation energy. It does not push the reaction forward with an energy input. It constrains geometry. The substrate enters the active site and its conformational freedom is drastically reduced — it is held in the orientation that places the reacting atoms in alignment with the catalytic residues, which are themselves positioned by the protein fold to deliver the correct geometry for bond formation or cleavage.

This explains two features of enzyme catalysis that are otherwise mysterious. First, enzymes are extremely fast: alignment does not need to occur by random conformational sampling, because the enzyme has pre-positioned the substrate in the aligned geometry. The alignment event is not rate-limiting because it is pre-satisfied by binding. Second, enzymes

are extremely selective: competing geometry routes are geometrically excluded by the active-site walls. A substrate molecule bound in the active site cannot adopt the conformation that would satisfy the alignment requirements of a competing pathway, because the protein structure physically prevents it. The active site is a geometry filter. It admits one alignment and blocks all others. The result is rate acceleration and selectivity simultaneously, both arising from the same geometry manipulation.

Organocatalysts, particularly proline-derived catalysts used in asymmetric enamine catalysis, work by a different geometry mechanism. These catalysts form a covalent intermediate with the substrate — an enamine — that introduces new stereocenters into the geometry of the reactive intermediate. The enamine geometry places one face of the pi bond in a specific orientation relative to the catalyst's chiral scaffold. The other face of the pi bond is blocked by the catalyst — specifically, by the carboxylic acid group of the proline-derived catalyst, which hydrogen-bonds to the incoming electrophile and directs its approach toward the exposed face. The catalyst has reshaped the geometry route so that only one face of the pi system is inside the tension window for electrophilic attack. The product is formed with high enantioselectivity not because one enantiomeric transition state is abstractly lower in energy, but because the catalyst geometry physically routes the electrophile to one face and blocks the other.

All three catalyst types share a common logic: they intervene in the geometry route between reactant and transition state. Lewis acids pre-distort the substrate geometry toward the transition state. Enzyme active sites pre-satisfy the alignment requirements. Organocatalysts form intermediates that filter which geometry routes are accessible. In each case, the energy change is real and measurable, but it is downstream of the geometry intervention. The geometry is the mechanism. The energy diagram is the record.

7. Pathway Failure Modes

A reaction pathway fails when the molecule cannot reach the geometry required by the transition state. This is always the proximate cause, regardless of how the failure presents experimentally. But geometry failure can occur for several distinct mechanical reasons, each with its own signature and its own fix. Conflating these failure modes — treating all pathway failures as vaguely "energetically unfavorable" — prevents systematic diagnosis and repair. There are five distinct mechanical failure modes, and they are each diagnosable from the product distribution or from the absence of product.

The first failure mode is the **tension-window miss**. The molecule's mechanical strain state is outside the window required for the geometry transition to occur. This can happen in either direction: the molecule is too relaxed and lacks pre-organization, or the molecule is too strained and its geometry has been deformed away from the transition-state approach geometry. A completely unactivated, unstrained ester at room temperature in a non-polar solvent may fail to react with a moderate nucleophile not because the reaction is thermodynamically unfavorable but because the molecule is fully relaxed and the thermal energy available is insufficient to drive the full geometric distortion from ground state to transition state without the geometric head start provided by pre-existing strain. The fix for a tension-window miss is to bring the molecule back into the window: raise temperature to increase thermal distortion, choose a solvent that interacts with the molecule in a way that introduces productive strain, or modify substituents to introduce directed strain toward the transition-state geometry.

The second failure mode is **alignment lock**. A conformational constraint — a rigid ring, a bulky substituent, a strong intramolecular hydrogen bond — holds the molecule in a geometry that cannot satisfy the alignment requirements of the desired transition state. The molecule has enough energy to react in principle, and it may be inside the appropriate tension window, but it is physically stuck in the wrong conformation. It cannot access the

geometry that would satisfy all four alignment constraints simultaneously. The E2 example from Section 3 is the archetype: the molecule is locked in a conformation where no beta hydrogen is anti-periplanar to the leaving group. The fix for alignment lock is to remove or weaken the conformational constraint: modify the ring, change a substituent that is enforcing the wrong conformation, choose a solvent that disrupts an intramolecular interaction that was enforcing the locked geometry, or redesign the substrate to allow the needed conformational excursion.

The third failure mode is **domain crowding**. Steric bulk around the reaction center physically prevents the approach of the reacting partner. The electron domains of the approaching molecule collide with the electron domains of the substrate before the alignment geometry is achieved. This is distinct from alignment lock: the substrate may be in a perfectly good conformation, but the trajectory the approaching reagent must follow is physically obstructed by substituents. Domain crowding is the failure mode operative in neopentyl SN₂, in hindered ketone additions, and in bulky electrophile aromatic substitution. The fix is not to use more reactive reagents — the electron-domain collision is a geometry constraint, not an energy constraint, and it does not yield to increased driving force. The fix is to redesign the approach geometry: use a smaller reagent, change the leaving group to one that is less obstructive, remove or modify substituents that are blocking the approach trajectory, or change the reaction site entirely.

The fourth failure mode is **geometry collapse**. During the reaction, an intermediate adopts a geometry that collapses to an unproductive configuration before the next step can proceed. A planar carbocation intermediate, for example, may have formed with the correct relative configuration for a subsequent stereospecific displacement, but the flat sp² geometry at the carbocation center means that the next nucleophile can attack from either face equally. The stereochemical information needed for the next step has been lost by geometry collapse. The fix requires keeping the intermediate in a productive geometry long enough for the next step to occur: add a nucleophile trap in the same reaction vessel to intercept the intermediate immediately after formation, accelerate the next step by

increasing the nucleophile concentration, or redesign the reaction sequence to avoid the problematic intermediate geometry entirely.

The fifth failure mode is **competitive drain**. A faster geometry route — a side-reaction pathway — drains the molecular population before the desired pathway can proceed. This is not a failure of the desired pathway's geometry per se. The desired pathway is geometrically accessible. But a competing pathway is more accessible — its alignment events are more frequent, its tension window is wider, or its alignment geometry is more easily reached from the current molecular geometry state. The result is that most of the molecular population is routed into the competing pathway before the desired pathway can capture it. The fix is the branch-window management discussed in Section 5: close the competing window by introducing geometry constraints that shut off the competing alignment, or open the desired window by removing constraints that were narrowing it.

8. Side-Reaction Geometry Drift

Side reactions are not random failures. They are geometry routes that open when the main pathway's tension window shifts or closes, or when a competing pathway's window widens. **Side-reaction geometry drift** describes the phenomenon where the molecular population drifts into the tension window for a competing pathway because conditions have shifted — temperature, solvent, concentration, or the structure of the molecule itself has changed — in a way that changes the relative accessibility of the competing geometry routes. Understanding side reactions as geometry drift events rather than as random or mysterious failures makes them diagnostic: the geometry of the side product tells you exactly which alignment window opened, and that information tells you what changed in the geometry-state distribution of the molecular population.

The most common side-reaction geometry drift in synthetic chemistry is **elimination when substitution was intended**. The molecule begins in a geometry state where both S_N2 and $E2$ alignment events are possible, but conditions favor S_N2 . As temperature rises

— because the reaction is being heated to speed it up — the conformational population shifts. Higher temperature populates gauche and eclipsed conformers more fully. Some of those conformers place a beta hydrogen in or near an anti-periplanar arrangement with the leaving group. The frequency of E2 alignment events increases. If the base is strong enough (and nucleophiles in polar aprotic solvents often are), the E2 pathway captures an increasing fraction of the molecular population. The product distribution shifts toward alkene. The geometry drift is real: the molecule's conformational population distribution has moved, and a previously narrow alignment window for E2 has widened. The fix is to lower the temperature, which returns the conformational population to geometry states that favor SN2 alignment, or to choose a nucleophile that is a poor base, narrowing the E2 alignment window regardless of conformation.

The second common side-reaction geometry drift is **over-reaction**. The product of the desired reaction is itself inside the tension window for a second reaction. The geometry of the product — its bond lengths, angles, electron-domain orientations — places it in a reactive conformation for the next electrophile, nucleophile, or base in the reaction mixture. The product geometry drifts back into a reactive tension window. This is particularly common in alkylation reactions where the product amine or alcohol is more nucleophilic or more reactive than the starting material, and in electrophilic aromatic substitution where the monosubstituted product is more activated toward further substitution than the starting arene. The solution is always to monitor the geometry state of the product and design conditions that push the product out of the reactive tension window before it can react again. Dilute conditions push intermolecular reactions out of the tension window by reducing collision frequency. Temperature reduction slows the conformational dynamics that would allow the product to satisfy the alignment requirements of the second reaction. Protecting-group strategies change the electron-domain geometry of the product, physically removing it from the reactive tension window.

The third side-reaction geometry drift involves **carbocation rearrangement**. A carbocation intermediate forms with a specific geometry: the positive center is at a trigonal

planar sp^2 carbon, with an empty p orbital perpendicular to the plane. This empty p orbital must be filled to complete the reaction. But before an external nucleophile arrives, the geometry of the adjacent bonds is checked. If an adjacent C-H or C-C bond is anti-periplanar to the empty p orbital — which means the bond is in the correct geometry to overlap with that orbital — electron density from that adjacent bond routes into the empty orbital. This is a 1,2-hydride or 1,2-alkyl shift. It is not a failure of mechanism. It is a geometry route that was open: the C-H or C-C sigma bond was geometrically aligned with the empty orbital, and electron-domain routing through that alignment was energetically favorable. The molecule took the route that its geometry made available.

The diagnostic utility of side-product geometry is substantial. The structure of a rearrangement product tells you the geometry of the intermediate at the moment of rearrangement: specifically, which C-H or C-C bond was anti-periplanar to the empty orbital when the shift occurred. Working backward from the rearranged product structure identifies the geometry of the carbocation intermediate and the alignment that was open. This is the same analysis applied to any side-reaction product. The geometry of the side product is not an embarrassment — it is information. It encodes the geometry-state history of the molecular population and tells you exactly which alignment window opened at which point in the reaction. A chemist who reads side products as geometry information has a diagnostic tool. A chemist who reads them as random failures has nothing.

Restoring the main pathway after geometry drift requires identifying the drift event and closing the window that opened. For temperature-driven elimination drift: lower temperature or switch to a harder, smaller nucleophile. For over-reaction drift: redesign product geometry through protecting groups or reaction stoichiometry. For rearrangement drift: redesign the intermediate to close the rearrangement alignment window, typically by reducing the stability of the carbocation (slower formation), accelerating the subsequent step, or designing conditions where an external nucleophile is present in high enough concentration to intercept the carbocation before the intramolecular rearrangement alignment event can occur.

9. Pathway Control via Tension and Geometry

Every decision a chemist makes about solvent, temperature, protecting groups, substituents, and catalysts is a geometry decision before it is anything else. The geometry-route framework makes this explicit and practical. Rather than asking "which conditions favor my desired product?" in a trial-and-error sense, the geometry-first approach asks a precise question: what geometry does the desired transition state require, and which conditions put the molecular population into the geometry state that best satisfies those requirements while closing the windows for competing pathways? This section lays out the five main control strategies, each grounded in a specific geometry mechanism.

Control Strategy 1: Tension Tuning via Solvent

Solvent choice is a tension-state decision. Polar aprotic solvents — dimethyl sulfoxide, acetone, dimethylformamide, acetonitrile — do not solvate anions effectively. A nucleophilic anion in a polar aprotic solvent retains its lone-pair geometry intact: the lone pair is not buried under a cage of hydrogen-bonding solvent molecules, and its orbital geometry is not distorted by tight solvation. This keeps the nucleophile's donor orbital in its natural geometry, available for alignment with the electrophile's acceptor orbital. Simultaneously, polar aprotic solvents stabilize charge-separated transition states through dipole interactions without preferentially solvating either the nucleophile or the developing negative charge on the leaving group. The tension window for the S_N2 pathway remains open, and the $E2$ window is not selectively widened. Polar aprotic conditions favor substitution over elimination for most nucleophile–substrate combinations.

Polar protic solvents — water, methanol, ethanol — solvate anions aggressively through hydrogen bonding. This wraps the nucleophile's lone pair in a cage of solvent that must be shed before the nucleophile can align with the electrophile. The effective geometry of the nucleophile is obscured: its lone pair is pre-occupied in hydrogen bonds that must break before alignment can occur. This raises the alignment-event frequency cost for S_N2 relative

to E2. Simultaneously, polar protic solvents stabilize developing charges through solvation, which widens the tension window for pathways that develop charge in the transition state — including E2, where the developing negative charge on the leaving group and the proton being abstracted are both stabilized by protic solvation. The product is a shift toward elimination. Choosing solvent is choosing which tension windows to open and which to narrow.

Control Strategy 2: Steric Geometry Sculpting

Substituents around the reaction center are geometry sculptures. A methyl group at a beta carbon does not simply add "steric bulk" in a vague sense. It changes the torsional strain landscape of the molecule in a specific way: it raises the energy of eclipsed conformers involving that methyl group, which shifts the equilibrium population of conformers away from those where the methyl is eclipsed and toward conformers where it is staggered. If the anti-periplanar E2 alignment geometry happens to place the beta methyl in a staggered position, the methyl group does not disfavor E2 by this mechanism. If it places the methyl in an eclipsed position, it does disfavor E2 because that conformer is less populated. Adding substituents is therefore a conformational population intervention, which is a direct geometry intervention in which alignment windows are open.

Control Strategy 3: Temperature as Conformational Population Dial

Temperature controls which conformational states are thermally accessible. At low temperature, molecules are trapped in or near their lowest-energy conformations. Conformations that are significantly strained — that require appreciable torsional or angle strain to reach — are not meaningfully populated. The geometry-state distribution of the molecular population is narrow, centered on the lowest-energy geometries. If the desired pathway's alignment event requires a low-energy conformation, low temperature favors that pathway. If the competing pathway's alignment event requires a strained, high-energy conformation, low temperature closes the competing window. High temperature populates strained conformations. Both windows that were closed at low temperature may open at high temperature. This is why running reactions cold often improves selectivity: you are

narrowing the geometry-state distribution of the molecular population and letting only the desired alignment geometry dominate.

Control Strategy 4: Directing Groups as Geometry Routers

A directing group in electrophilic aromatic substitution is not usefully described as making certain ring carbons "more electron-rich" in an abstract sense. That description is accurate at one level but obscures the geometry mechanism. A directing group orients the incoming electrophile's approach by establishing which orbital geometry of the pi system is most accessible and most aligned for electrophilic attack. An ortho/para director — a hydroxyl group, for instance — has its lone pair in conjugation with the ring pi system at the ortho and para positions. The pi electron density at those positions is higher, but more importantly, the geometry of the pi system at those positions is such that an incoming electrophile's empty orbital can overlap with the pi system most efficiently. The director is routing the approach geometry of the electrophile. Meta directors reduce pi density at ortho and para positions through electron withdrawal, but again, the practical effect is on which geometry of approach is most favorable. The directing effect is a geometry routing effect dressed in the language of electron density.

Control Strategy 5: Protecting Groups as Geometry Reservoirs

A protecting group does not merely block a functional group by covering it up. It changes the electron-domain geometry of the protected site in a way that removes that site from its reactive tension window. An alcohol converted to a silyl ether has its oxygen lone pairs now sharing electron density with the silicon atom. The oxygen's lone-pair geometry changes: instead of two freely available lone pairs pointing away from the C–O bond in classic sp^3 orientations, the lone pairs are partly committed to the O–Si bond, which is longer and weaker than the O–H bond. The electron-domain geometry of the oxygen is different. The protected oxygen is no longer in the tension window for the reactions that the unprotected alcohol would undergo, not merely because the oxygen is "blocked" but because its geometry has been changed. The protection strategy is a targeted geometry intervention

that shifts a specific site out of its reactive tension window while leaving other sites in theirs.

10. Restoring Geometry Routes

When a reaction fails to give the desired product, the instinctive response is to change the reagent. Use a stronger nucleophile. Use a different base. Add more catalyst. These changes sometimes work, but they work by accident when they happen to fix the underlying geometry problem, and they fail predictably when the underlying geometry problem is not addressed. The geometry-first instinct is different: ask which geometry route broke down and why, then apply a fix matched to the specific failure mode. This procedure is systematic because geometry is systematic. The molecule is doing what its geometry tells it to do. Understanding the geometry tells you what to fix.

Step 1 is diagnosis. Identify which failure mode from Section 7 has occurred. The diagnosis comes from the product distribution, or from the absence of product, and from careful attention to what is present that should not be. Tension-window miss leaves no product: the reaction simply does not proceed, yields are negligible, and varying reagent reactivity does not help much. Alignment lock produces a product from a competing pathway that does satisfy its alignment requirements, or produces no product at all if no competing pathway is accessible. Domain crowding produces little or no product when a bulky reagent is used, but product appears immediately when a smaller reagent is tried. Geometry collapse produces product with incorrect or scrambled stereochemistry — what should be a stereospecific reaction gives a racemate or an unexpected diastereomer. Competitive drain produces the wrong product in high yield — the competing pathway is working perfectly, and the desired pathway is starved.

Step 2 is geometry mapping. Identify precisely which bond angle, dihedral angle, or orbital orientation is not being achieved. This step requires building a model — physical or computational — of the required transition-state geometry and checking whether the

substrate as designed can thread through it. For computational work, this means optimizing the transition state and checking whether the geometry is accessible from the reactant conformation. For physical model work, this means building the reactant in its ground-state geometry and testing whether the transition-state conformation can be reached without unreasonable strain. The goal of this step is a specific geometric diagnosis: the H-C-C-LG dihedral cannot reach 180° because of this ring constraint; the nucleophile approach trajectory is blocked by this methyl group at this position; the carbocation intermediate flattens before the nucleophile can arrive. Specific, spatial, three-dimensional diagnoses.

Step 3 is strategy selection matched to failure mode. For a tension-window miss, adjust the strain state of the molecule: change the solvent to one that interacts with the reacting site in a strain-producing way, raise or lower temperature to shift the molecular tension distribution, add a Lewis acid that pre-strains the substrate geometry toward the transition state, or modify substituents to introduce directed strain. For an alignment lock, remove the conformational constraint: redesign the substrate to break the ring or remove the substituent enforcing the wrong conformation, change the solvent to disrupt an intramolecular interaction that was locking the geometry, or use a different reaction sequence that avoids passing through the locked geometry state. For domain crowding, redesign the approach geometry: switch to a smaller nucleophile or electrophile, choose a leaving group that creates less steric obstruction, remove or replace the substituents that are blocking the approach trajectory, or redesign the protecting-group strategy to clear the path to the reaction center. For geometry collapse, accelerate the subsequent step or add a trap: use a higher concentration of the next reagent, run the steps without isolation of the intermediate, or choose conditions where the subsequent step is faster than the collapse event. For competitive drain, close the competing window through the strategies of Section 9: solvent choice, substituent design, temperature management.

Step 4 is restoration verification through geometry prediction. After applying the fix, predict the geometry state of the substrate under the new conditions before running the

reaction again. Check whether the modified substrate now satisfies the alignment requirements of the desired transition state. If the geometry analysis says the route is restored — the dihedral can reach 180° , the approach trajectory is clear, the intermediate no longer collapses before the next step — then proceed. If the geometry analysis says the route is still blocked, the fix was wrong or incomplete. Return to Step 2, refine the geometric diagnosis, and select a better-matched strategy.

This four-step procedure is not a theoretical exercise. It is what experienced synthetic chemists do implicitly when a reaction fails and they redesign around the failure. Making the procedure explicit and geometry-grounded speeds up the design cycle and prevents the common error of applying the same fix repeatedly to a problem that requires a different kind of fix. A failed reaction is a geometry diagnosis problem. It has a specific cause, a specific geometric signature, and a specific class of fixes matched to that cause. Mechanism is not guesswork. The molecule is doing exactly what its geometry tells it to do, every time, without exception. Once you understand the geometry, you understand the machine.

Conclusion

Ten sections, one argument: chemistry is mechanical. Molecules move through geometry space, and the pathway they take is determined by which geometry routes are open. The reaction coordinate diagram is not the cause of anything — it is a downstream summary of geometry decisions the molecule has already made, atom by atom, bond angle by bond angle, electron domain by electron domain.

Section 1 established that pathways are sequences of geometry states, not abstract energy slopes. The S_N2 reaction is a geometry threading problem. Every reaction pathway is. Section 2 introduced the tension window: the range of mechanical strain within which a geometry transition is accessible. A pathway is gated by tension, not by activation energy in isolation. Activation energy is the cost of reaching the required geometry; the tension window defines whether that cost can be paid from the molecule's current mechanical

state. Section 3 established that geometry alignment is the actual gate: the simultaneous satisfaction of multiple spatial constraints, all of which must be met at once for a bond-forming or bond-breaking event to proceed. E2 anti-periplanar requirements are the prototype, but every reaction has its alignment requirements, and none of them negotiate.

Section 4 explained electron-domain routing as the physical basis of curved-arrow notation. Electrons move through geometry pathways. Conjugation, hyperconjugation, and lone-pair donation are all geometry routing phenomena. Where electrons go is determined by which orbital geometry routes are open. Section 5 analyzed selectivity as branch-window competition: multiple pathways fight over the same molecular population, and product distribution reflects the relative frequency of alignment events for each competing route. Section 6 showed that catalysts are geometry-reshaping tools, not abstract energy-lowering tools. Lewis acids pre-distort substrate geometry. Enzyme active sites pre-satisfy alignment requirements. Organocatalysts filter which geometry routes are accessible. The energy change is the consequence of the geometry change.

Section 7 cataloged five mechanical failure modes — tension-window miss, alignment lock, domain crowding, geometry collapse, and competitive drain — each with its own signature and its own class of fixes. Section 8 reframed side reactions as geometry drift events: the molecular population drifts into a competing pathway's tension window when conditions shift. Side-product geometry is diagnostic, not merely embarrassing. Section 9 synthesized the practical control strategies: solvent as tension-state selector, substituents as geometry sculptors, temperature as conformational population dial, directing groups as orbital geometry routers, and protecting groups as geometry reservoirs. Section 10 provided the systematic procedure for diagnosing and restoring failed geometry routes in four steps: diagnose the failure mode, map the missing geometry, select a matched restoration strategy, verify by geometry prediction.

The energy diagram is not wrong. It correctly summarizes the cost profile of the reaction. But it is not explanatory. It does not tell you why the molecule takes the pathway it takes, why a competing pathway is preferred, why the reaction failed, or what to change to fix it.

The geometry does. Understanding the geometry gives a chemist actual control over the machine. The energy diagram records the outcome; the geometry makes the decision.

Implications

The geometry-route framework changes how you design reactions. You stop asking "what is the activation energy?" first and start asking "what geometry does the transition state require, and is my molecule currently capable of threading through that geometry?" This shifts design work upstream. Substituent choice, solvent choice, temperature selection, protecting-group strategy, and catalyst design are all geometry decisions before they are anything else. Every one of these variables acts by changing the geometry-state distribution of the molecular population, opening some tension windows and closing others, enabling some alignment events and blocking others. Treating them as geometry decisions — rather than as vague optimization knobs — makes their effects predictable and their failures diagnosable.

A failed reaction is a geometry diagnosis problem, not a mystery. It has a specific mechanical cause, a specific geometric signature, and a specific class of fixes matched to that cause. And it changes how mechanism should be taught. Arrow-pushing is not a bookkeeping exercise. It is a geometry routing map that tells you where electron density is moving and what spatial constraints must be satisfied for it to move there. Teach it that way — with explicit attention to orbital orientation, conformational requirements, and alignment constraints — and students will understand why the arrows go where they go. They will stop memorizing patterns and start reading machines.